

# **SDS 385: Exercise 01**

August 27, 2016

**Jennifer Starling**



## Exercises 1: Preliminaries

### Linear regression

Consider the simple linear regression model

$$y = X\beta + e,$$

where  $y = (y_1, \dots, y_N)$  is an  $N$ -vector of responses,  $X$  is an  $N \times P$  matrix of features whose  $i$ th row is  $x_i$ , and  $e$  is a vector of model residuals. The goal is to estimate  $\beta$ , the unknown  $P$ -vector of regression coefficients.

Let's say you trust the precision of some observations more than others, and therefore decide to estimate  $\beta$  by the principle of weighted least squares (WLS):

$$\hat{\beta} = \arg \min_{\beta \in \mathbb{R}^P} \sum_{i=1}^N \frac{w_i}{2} (y_i - x_i^T \beta)^2,$$

where  $w_i$  is the weight for observation  $i$ . (Higher weight means more influence on the answer; the factor of  $1/2$  is just for convenience, as you'll see later.)

- (A) Rewrite the WLS objective function<sup>1</sup> above in terms of vectors and matrices, and show that  $\hat{\beta}$  is the solution to the following linear system of  $P$  equations in  $P$  unknowns:

$$(X^T W X) \hat{\beta} = X^T W y,$$

where  $W$  is the diagonal matrix of weights.

- (B) One way to calculate  $\hat{\beta}$  is to: (1) recognize that, trivially, the solution to the above linear system must satisfy  $\hat{\beta} = (X^T W X)^{-1} X^T W y$ ; and (2) to calculate this directly, i.e. by inverting  $X^T W X$ . Let's call this the "inversion method" for calculating the WLS solution.

Numerically speaking, is the inversion method the fastest and most stable way to actually solve the above linear system? Do some independent sleuthing on this question. Summarize what you find, and provide pseudo-code for at least one alternate method based on matrix factorizations—call it "your method" for short. (Note: our linear system is not a special flower; whatever you discover about general linear systems should apply here.)

- (C) Code up functions that implement both the inversion method and your method for an arbitrary  $X$ ,  $y$ , and weights  $W$ . Obviously you

<sup>1</sup> That is, the thing to be minimized.

shouldn't write your own linear algebra routines for doing things like multiplying or decomposing matrices, but don't use a direct model-fitting function like R's `lm` either. Your actual code should look a lot like the pseudo-code you wrote for the previous part. Note: be attentive to how you multiply a matrix by a diagonal matrix, or you'll waste a lot of time multiplying stuff by zero.

Now simulate some silly data from the linear model for a range of values of  $N$  and  $P$ . (Feel free to assume that the weights  $w_i$  are all 1.) It doesn't matter how you do this—e.g. everything can be Gaussian if you want. (We're not concerned with statistical principles yet, just with algorithms, and using least squares is a pretty terrible idea for enormous linear models, anyway.) Just make sure that you explore values of  $P$  up into the thousands, and that  $N > P$ . Benchmark the performance of the inversion solver and your solver across a range of scenarios. (In R, a simple library for this purpose is `microbenchmark`.)

- (D) Now what happens if  $X$  is a highly sparse matrix, in the sense that most entries are zero? Ideally we'd realize some savings by not doing a whole bunch of needless multiplication by zero in our code.

It's easy to simulate an  $X$  matrix that looks like this. A quick-and-dirty way is to simulate a mask of zeros and ones (but mostly zeros), and then do pointwise multiplication with your original feature matrix. For example:

```
N = 2000
P = 500
X = matrix(rnorm(N*P), nrow=N)
mask = matrix(rbinom(N*P,1,0.05), nrow=N)
X = mask*X
X[1:10, 1:10] # quick visual check
```

Again assume that the weights  $w_i$  are all 1. Repeat the previous benchmarking exercise with this new recipe for simulating a sparse  $X$ , except add another solver to the mix: one that can solve a linear system  $Ax = b$  in a way that exploits the sparsity of  $A$ . To do this, you'll need to actually represent the feature matrix  $X$  in a sparse format, and then call the appropriate routines for that format. (Again, do some sleuthing; in R, the `Matrix` library has data structures and functions that can do this; SciPy will have an equivalent.)

Benchmark the inversion method, your method, and the sparse

method across some different scenarios (including different sparsity levels in  $X$ , e.g. 5% dense in my code above).

### Generalized linear models

As an archetypal case of a GLM, we'll consider the binomial logistic regression model:  $y_i \sim \text{Binomial}(m_i, w_i)$ , where  $y_i$  is an integer number of "successes,"  $m_i$  is the number of trials for the  $i$ th case, and the success probability  $w_i$  is a regression on a feature vector  $x_i$  given by the inverse logit transform:

$$w_i = \frac{1}{1 + \exp\{-x_i^T \beta\}}.$$

We want to estimate  $\beta$  by the principle of maximum likelihood. Note: for binary logistic regression,  $m_i = 1$  and  $y_i$  is either 0 or 1.

As an aside, if you have a favorite data set or problem that involves a different GLM—say, a Poisson regression for count data—then feel free to work with that model instead throughout this entire section. The fact that we're working with a logistic regression isn't essential here; any GLM will do.

(A) Start by writing out the negative log likelihood,

$$l(\beta) = -\log \left\{ \prod_{i=1}^N p(y_i | \beta) \right\}.$$

Simplify your expression as much as possible. This is the thing we want to minimize to compute the MLE. (By longstanding convention, we phrase optimization problems as minimization problems.)

Derive the gradient of this expression,  $\nabla l(\beta)$ . Note: your gradient will be a sum of terms  $l_i(\beta)$ , and it's OK to use the shorthand

$$w_i(\beta) = \frac{1}{1 + \exp\{-x_i^T \beta\}}$$

in your expression.

(B) Read up on the method of steepest descent, i.e. gradient descent, in Nocedal and Wright (see course website). Write your own function that will fit a logistic regression model by gradient descent. Grab the data "wdbc.csv" from the course website, or obtain some other real data that interests you, and test it out. The WDBC file has information on 569 breast-cancer patients from a study done in Wisconsin. The first column is a patient ID, the second column is a

classification of a breast cell (Malignant or Benign), and the next 30 columns are measurements computed from a digitized image of the cell nucleus. These are things like radius, smoothness, etc. For this problem, use the first 10 features for  $X$ , i.e. columns 3-12 of the file. If you use all 30 features you'll run into trouble.

Some notes here:

1. You can handle the intercept/offset term by either adding a column of 1's to the feature matrix  $X$ , or by explicitly introducing an intercept into the linear predictor and handling the intercept and regression coefficients separately, i.e.

$$w_i(\beta) = \frac{1}{1 + \exp\{-(\alpha + x_i^T \beta)\}}.$$

2. I strongly recommend that you write a self-contained function that, for given values of  $\beta$ ,  $y$ ,  $X$ , and sample sizes  $m_i$  (which for the WDBC data are all 1), will calculate the gradient of  $l(\beta)$ . Your gradient-descent optimizer will then call this function. Modular code is reusable code.
3. Make sure that, at every iteration of gradient descent, you compute and store the current value of the log likelihood, so that you can track and plot the convergence of the algorithm.
4. Be sensitive to the numerical consequences of an estimated success probability that is either very near 0, or very near 1.
5. Finally, you can be as clever as you want about the gradient-descent step size. Small step sizes will be more robust but slower; larger step sizes can be faster but may overshoot and diverge; step sizes based on line search (Chapter 3 of Nocedal and Wright) are cool but involve some extra work.

- (C) Now consider a point  $\beta_0 \in \mathcal{R}^p$ , which serves as an intermediate guess for our vector of regression coefficients. Show that the second-order Taylor approximation of  $l(\beta)$ , around the point  $\beta_0$ , takes the form

$$q(\beta; \beta_0) = \frac{1}{2}(z - X\beta)^T W(z - X\beta) + c,$$

where  $z$  is a vector of "working responses" and  $W$  is a diagonal matrix of "working weights," and  $c$  is a constant that doesn't involve  $\beta$ . Give explicit expressions for the diagonal elements  $W_{ii}$  and for  $z_i$  (which will necessarily involve the point  $\beta_0$ , around which you're doing the expansion).<sup>2</sup>

<sup>2</sup> Remember the trick of completing the square, e.g. <https://justindomke.wordpress.com/completing-the-square-in-n-dimensions/>.

- (D) Read up on Newton's method in Nocedal and Wright, Chapter 2. Implement it for the logit model and test it out on the same data set you just used to test out gradient descent.<sup>3</sup> Note: while you could do line search, there is a "natural" step size of 1 in Newton's method.
- (E) Reflect broadly on the tradeoffs inherent in the decision of whether to use gradient descent or Newton's method for solving a logistic-regression problem.

<sup>3</sup> You should be able to use your own solver for linear systems from the first section.

## Linear Regression

### Part A

The WLS objective function, rewritten in matrix form, is:

$$\hat{\beta} = \arg \min_{\beta \in R^p} \frac{1}{2} (Y - X'\beta)'W(Y - X'\beta) = \arg \min_{\beta \in R^p} \frac{1}{2} (Y' - X'\beta')W(Y - X'\beta)$$

To satisfy the 'arg min' part of the expression, take the derivative of  $\hat{\beta}$  with respect to  $\beta$ , set equal to zero, and solve as follows.

$$\frac{\delta}{\delta\beta} [\frac{1}{2} (Y' - X'\beta')W(Y - X'\beta)] = (\frac{1}{2}) \frac{\delta}{\delta\beta} [(Y'WY - 2Y'WX\beta + \beta'XWX'\beta)] = 0$$

The derivatives of each term are as follows.

- (a)  $\frac{\delta}{\delta\beta} [Y'WY] = 0$  since this term is constant with respect to  $\beta$ .
- (b)  $\frac{\delta}{\delta\beta} [-2Y'WX\beta] = -2Y'WX$ , since derivative has form  $\frac{\delta}{\delta\beta} c\beta = c$ , where  $c = -2Y'WX$ . Then since  $W$  diagonal (and so symmetric) and  $X$  and  $Y$  vectors,  $-2Y'WX = -2X'WY$ .
- (c)  $\frac{\delta}{\delta\beta} [\beta'XWX'\beta] = 2XWX'\beta$ , since derivative has quadratic form  $\frac{\delta}{\delta\beta} [\beta'V\beta] = (V + V')\beta$ , where  $V = XWX'$ . When  $V$  symmetric, this further simplifies to  $2V\beta$ .

Then the derivative, subbing in  $\hat{\beta}$  for  $\beta$ , is

$$\frac{1}{2} [-2X'WY + 2XWX'\hat{\beta}] = 0 \rightarrow XWX'\hat{\beta} = X'WY \rightarrow \hat{\beta} = (XWX')^{-1}X'WY.$$

Therefore  $\hat{\beta} = (XWX')^{-1}X'WY$ .

To show that  $\hat{\beta} = (XWX')^{-1}X'WY$  is the solution to the linear system:

$$(X'WX)\hat{\beta} = (X'WX)(X'WX)^{-1}X'WY = IX'WY = X'WY.$$



**Part B**

Numerically speaking, I do not believe that inversion is the fastest and most stable way to solve the linear system. There are several matrix factorization methods which provide more stability and are computationally efficient compared to inversion. Inverting a matrix directly is computationally intensive, especially as  $N$  and  $P$  become large.

Some of the methods I discovered for solving linear equations of form  $Ax=B$  without inverting the  $A$  matrix directly are: LU Decomposition, Gaussian elimination, Cholesky decomposition, QR decomposition, RRQR factorization, and the conjugate gradient method. There was not a strict consensus as to which method is universally superior; the key to know the characteristics of the matrix  $A$ , so that you can choose an optimal method. Different characteristics lend themselves to different methods.

- (a) Cholesky performs well for Hermitian matrices (symmetric positive definite).
- (b) LU performs well when  $A$  is sparse, and  $A$  is only required to be square.
- (c) Conjugate gradient requires  $A$  to be symmetric positive definite, but is a good iterative algorithm for scenarios where  $A$  is sparse and too large to be inverted directly or for Cholesky.

My method will be the Cholesky decomposition. My pseudo-code is as follows.

---

Goal: Solve  $Ax = b$  where  $A = X'WX$ ,  $b = X'WY$ , and  $x$  is the vector of  $\beta$  coefficient estimates.

Function inputs:

- (a)  $X$ , an  $N \times P$  matrix
- (b)  $W$ , a diagonal matrix of weights
- (c)  $Y$ , an  $N \times 1$  vector of responses

Function outputs:

- (a)  $\hat{\beta}$ , a vector of coefficient estimates

Code Steps:

- (1) Set  $A = X'WX$
- (2) Set  $b = X'WY$
- (3) Set  $R$  = Cholesky decomposition of  $A$ . ( $R$  gives the  $R$  (upper) instead of  $L$  (lower).)

Now we have  $R'R = A$ .

- (4) Solve  $R'z = b$  for  $z$  by finding  $z = (R')^{-1}b$
  - (4) Solve  $Rx = z$  by finding  $x = R^{-1}z$
  - (4) Return  $\hat{\beta}$  estimate as  $\hat{\beta} = x$
- 

I also included the LU decomposition, which is solved by similar steps.

## Part C

The R code for implementing and benchmarking the functions is as follows.

```

### SDS 385 - Exercises 01 - Part A
#This code compares various matrix decomposition
#methods to the inversion method, and benchmarks
#performance of the Cholesky and LU methods versus
5 #inversion at various sample/parameter sizes,
#and various sparsity levels of the X matrix.

#Jennifer Starling
#22 August 2016

10 library(Matrix) #For matrix decomposition.
library(microbenchmark) #For benchmarking

### PART C:

15 #-----
#Inversion Method function:
# Inputs: X = vector of x values, Y = vector of y values,
#         W = diag matrix of weights.
20 # Outputs: B = beta-hat vector; the WLS solution for
#           estimating the beta vector of coefficients.

# Matrix requirements:
# 1. Length X = Length Y = Dim(W)
25 # 2. t(X) %*% W %*% X must be invertible

inv_method <- function(X,W,y){
  B_hat <- solve(t(X) %*% W %*% X) %*% t(X) %*% W %*% y
  return(B_hat)
30 }

#-----
#LU Decomposition to solve linear system Ax=b.
# Inputs: X = vector of x values, Y = vector of y values,
35 #         W = diag matrix of weights.
# Outputs: B_hat_LU, an estimate of the 'x' in Ax=b.

lu_method <- function(X,W,y){

40   #Solves linear system Ax=b.
   #Since we have (X'WX)B=X'Wy, B (beta) acts as x, with A and b as follows.

   #Finding B (beta_hat) in equation
   A = (t(X) * diag(W)) %*% X #Efficient way of A = t(X) %*% W %*% X as W diag.
45   #Avoids mult by 0's.
   b = (t(X) * diag(W)) %*% y #b'Wy

   #Obtain LU matrix decomposition of A.
   decomp <- lu(A) #Calculates matrix decomposition object..
50   L <- expand(decomp)$L #Upper triangular matrix

```

```

    U <- expand(decomp)$U    #Lower triangular lower triangular matrix
    #Note: Uses partial pivoting. $P shows pivot matrix.

    #Now we replace Ax=b with LUX=b.
55  #Introduce Ld=b, giving us two linear equation systems: Ld=b and Ux=d.
    #So we will solve in two steps.

    #1. Solve Ld=b for d. This is d=inv(L)b
    d <- solve(L) %*% b
60

    #2. Substitutde d into Ux=d to solve for x. This is x = inv(U)d. (x=beta_hat)
    B_hat_LU <- solve(U) %*% d

    return(B_hat_LU)        #Returns function output.
65 }

#-----
#Cholesky decomposition function:
# Inputs: X = vector of x values, Y = vector of y values,
70 # W = diag matrix of weights.
# Outputs: B = beta-hat vector; the WLS solution for
# estimating the beta vector of coefficients.

cholesky_method <- function(X,W,y){
75  #Solves linear system Ax=b.
    #Since we have (X'WX)B=X'Wy, B (beta) acts as x, with A and b as follows.

    #Finding B (beta_hat) in equation
    A = (t(X) * diag(W)) %*% X #Efficient way of A = t(X) %*% W %*% X as W diag.
80  #Avoids mult by 0's.
    b = (t(X) * diag(W)) %*% y #b'Wy
    R <- chol(A) #Find right/upper cholesky decomposition of A.

    #Now we have R'R=A.
85

    #1. Solve R'z=b for z. This is z = inv(R')b.
    z = solve(t(R)) %*% b

    #2. Solve Rx=z for x. This is x = inv(R)z. (x = beta_hat)
90  B_hat_chol <- solve(R) %*% z

    return(B_hat_chol)
}

#-----
95 #BENCHMARKING:

#Simulate data from the linear model for a range of values of N and P.
#(Assume weights are all 1, data are gaussian.)
100 #Carry out performance testing of the two methods.

library(microbenchmark)

```

```

N <- c(10,100,500,1000)
105 P <- N/2 #Setting up so that N>P. This is an arbitrary choice.

perf_results <- list()

110 for (i in 1:length(N)){

    n <- N[i]
    p <- P[i]

    print(n)

115 #Set up matrices of size N, P parameters: (dummy data)
    X <- matrix(rnorm(n*p), nrow=n, ncol=p)
    y <- rnorm(n)
    W <- diag(1, nrow=n)

120 #Perform benchmarking:
    perf_results[[i]] <- microbenchmark(
        inv_method(X,W,y),
        lu_method(X,W,y),
125 cholesky_method(X,W,y), unit='ms'
    )
}

names(perf_results) <- (c('N=10,P=5', 'N=100,P=50', 'N=500,P=250', 'N=1500,P=500'))
130 perf_results #Display benchmarking results.

```

The performance benchmarking results are as follows. The inverse method was fastest for very small N and P values, but as N and P increased, LU and Cholesky performed more quickly than inverse. LU was the fastest method of the three.

```

> perf_results
$`N=10,P=5`
Unit: milliseconds
      expr      min       lq     mean  median      uq      max neval cld
inv_method(X, W, y) 0.038848 0.0442330 0.1654871 0.0482625 0.053931 10.973424 100 a
lu_method(X, W, y) 0.097958 0.1094925 0.1618572 0.1168010 0.134071 2.297167 100 a
cholesky_method(X, W, y) 0.089010 0.1019180 0.1291046 0.1108950 0.122683 0.445684 100 a

$`N=100,P=50`
Unit: milliseconds
      expr      min       lq     mean  median      uq      max neval cld
inv_method(X, W, y) 1.341858 1.3730145 1.5708283 1.464907 1.6235910 3.893104 100 c
lu_method(X, W, y) 0.389776 0.4115185 0.5163842 0.443505 0.5392765 2.367135 100 a
cholesky_method(X, W, y) 0.621448 0.6396920 0.8082234 0.709129 0.7825670 3.414895 100 b

$`N=500,P=250`
Unit: milliseconds
      expr      min       lq     mean  median      uq      max neval cld
inv_method(X, W, y) 138.48695 143.65306 147.27360 145.88347 148.95680 236.5789 100 c
lu_method(X, W, y) 27.08385 28.66972 31.97459 29.76773 31.06965 123.7989 100 a
cholesky_method(X, W, y) 52.35322 54.36095 56.99425 55.57933 57.59481 143.0048 100 b

$`N=1500,P=500`
Unit: milliseconds
      expr      min       lq     mean  median      uq      max neval cld
inv_method(X, W, y) 1159.7687 1181.4272 1204.1996 1192.2028 1210.0858 1318.4425 100 c
lu_method(X, W, y) 217.5088 226.2489 236.3857 231.3792 237.0129 366.1290 100 a
cholesky_method(X, W, y) 425.7139 435.8642 450.1871 441.6242 450.4200 561.5078 100 b

```

Figure 1: Performance benchmarking with dense X matrix

## Part D

Since both LU and Cholesky are good for sparse matrices, I benchmarked both of these methods against the inverse method for a sparse matrix X.

I performed two types of benchmarking:

- (a) Benchmarking various N and P at 10% sparse.
- (b) Benchmarking various sparsity levels (5%, 10%, 20%, 50%).

For benchmarking at various N and P levels, results were similar to the results above. Inverse was superior for small N and P, and as N and P increased, Cholesky and LU were superior. LU again performed the most efficiently.

```
> perf_results #Display benchmarking results.
$`N=10,P=5`
Unit: microseconds
      expr      min       lq     mean    median      uq     max neval cld
inv_method(X, W, y) 195.270 230.6890 250.7306 246.817 263.0400 436.975 100 ab
lu_method(X, W, y) 142.961 185.6895 233.2497 208.937 236.2485 794.388 100 a
cholesky_method(X, W, y) 184.360 223.1060 269.7928 244.267 278.2675 708.212 100 b

$`N=100,P=50`
Unit: microseconds
      expr      min       lq     mean    median      uq     max neval cld
inv_method(X, W, y) 1340.657 1362.3990 1461.4733 1452.5095 1491.339 2091.966 100 c
lu_method(X, W, y) 393.144 429.8705 557.8016 506.2400 530.219 5921.085 100 a
cholesky_method(X, W, y) 631.044 655.0500 754.6730 739.4485 773.096 2299.500 100 b

$`N=500,P=250`
Unit: milliseconds
      expr      min       lq     mean    median      uq     max neval cld
inv_method(X, W, y) 139.28952 147.35422 152.22056 150.50908 153.99406 254.34373 100 c
lu_method(X, W, y) 27.87053 30.68015 32.11990 31.90189 33.13519 39.19118 100 a
cholesky_method(X, W, y) 52.88206 56.24804 58.58854 58.02186 60.73400 71.86568 100 b

$`N=1000,P=500`
Unit: milliseconds
      expr      min       lq     mean    median      uq     max neval cld
inv_method(X, W, y) 1140.2651 1206.4823 1407.2283 1266.9112 1527.1099 2560.0583 100 c
lu_method(X, W, y) 220.1672 234.8487 299.8574 251.6809 300.5480 782.2592 100 a
cholesky_method(X, W, y) 419.3958 446.5509 497.6972 469.8177 524.3121 883.4575 100 b
```

Figure 2: Performance benchmarking with sparse X matrix

For benchmarking at various sparsity levels (with N=100, P=50 for all levels), the LU method performed most efficiently again, followed by Cholesky. For all methods, performance slowed as the matrix became less sparse.

```
> perf_results
$`5%`
Unit: microseconds
      expr      min       lq     mean    median      uq     max neval cld
inv_method(X, W, y) 1339.781 1359.7520 1469.3973 1383.2195 1479.3050 2624.806 100 c
lu_method(X, W, y) 393.088 421.8870 464.1313 435.1490 460.1350 1081.836 100 a
cholesky_method(X, W, y) 654.567 679.9915 803.9488 693.1765 744.2205 6065.418 100 b

$`10%`
Unit: microseconds
      expr      min       lq     mean    median      uq     max neval cld
inv_method(X, W, y) 1341.014 1358.3260 1452.5574 1388.2360 1483.1140 2098.072 100 c
lu_method(X, W, y) 388.911 411.5860 465.5492 430.5955 463.8015 1295.116 100 a
cholesky_method(X, W, y) 618.995 637.5325 744.7535 648.2635 701.9220 6531.208 100 b

$`25%`
Unit: microseconds
      expr      min       lq     mean    median      uq     max neval cld
inv_method(X, W, y) 1341.192 1361.8625 2424.2772 1384.4855 1446.653 101512.188 100 b
lu_method(X, W, y) 393.414 410.0670 474.5369 425.5910 473.378 1209.219 100 a
cholesky_method(X, W, y) 622.194 638.1095 707.9242 647.6305 694.926 1448.988 100 ab

$`50%`
Unit: microseconds
      expr      min       lq     mean    median      uq     max neval cld
inv_method(X, W, y) 1342.665 1357.659 1432.2028 1368.9885 1467.3945 2316.136 100 c
lu_method(X, W, y) 389.675 414.527 511.9470 432.1205 466.6205 6331.394 100 a
cholesky_method(X, W, y) 619.193 638.728 685.6965 657.8495 705.1420 1335.383 100 b
```

Figure 3: Performance benchmarking with varying sparsity levels of X

## Generalized Linear Regression

### Part A

The negative log-likelihood function is simplified as follows.

$$\begin{aligned} l(\beta) &= -\log\{\prod_{i=1}^N p(y_i|\beta)\} = -\log\{\prod_{i=1}^N \binom{m_i}{y_i} w_i^{y_i} (1-w_i)^{(m_i-y_i)}\} \\ &= -\sum_{i=1}^N \log\{w_i^{y_i} (1-w_i)^{(m_i-y_i)}\} = -\sum_{i=1}^N \{y_i \log(w_i) + (m_i - y_i) \log(1-w_i)\}; w_i = \frac{1}{1+\exp(x'_i \beta)} \end{aligned}$$

The gradient is found by taking the derivative of  $l(\beta)$  with respect to  $\beta$ .

First, the derivative of  $w_i$  with respect to  $\beta$  will be useful:

$$\frac{\delta w_i}{\delta \beta} = -(1 + \exp(-x'_i \beta))^{-2} \cdot \frac{\delta}{\delta \beta} (\exp(-x'_i \beta)) = \frac{-\exp(-x'_i \beta)(-x'_i)}{(1 + \exp(-x'_i \beta))^2} = \frac{x'_i \exp(-x'_i \beta)}{(1 + \exp(-x'_i \beta))^2}$$

Then find the gradient:

$$\begin{aligned} \frac{\delta l(\beta)}{\delta \beta} &= -\sum_{i=1}^N \{y_i \frac{1}{w_i} (\frac{\delta w_i}{\delta \beta}) + (m_i - y_i) \frac{1}{1-w_i} (-\frac{\delta w_i}{\delta \beta})\} = -\sum_{i=1}^N \{y_i \frac{1}{w_i} (\frac{x_i \exp(-x'_i \beta)}{1 + \exp(-x'_i \beta)}) + (m_i - y_i) \frac{1}{1-w_i} (\frac{-x_i \exp(-x'_i \beta)}{(1 + \exp(-x'_i \beta))^2})\} \\ &= -\sum_{i=1}^N \{y_i \frac{1}{w_i} (w_i^2 x_i \exp(-x'_i \beta)) + (m_i - y_i) \frac{1}{1-w_i} (w_i^2 x_i \exp(-x'_i \beta))\} \\ &= -\sum_{i=1}^N \{y_i w_i \exp(-x'_i \beta) - (m_i - y_i) x_i w_i\} = -\sum_{i=1}^N \{(y_i w_i \exp(-x'_i \beta) - (m_i - y_i) w_i) x_i\} \end{aligned}$$

And since  $w_i = \frac{1}{1+\exp(-x'_i \beta)} \rightarrow \exp(-x'_i \beta) = \frac{1}{w_i} - 1$ , we can simplify further:

$$\nabla l(\beta) = -\sum_{i=1}^N \{(y_i w_i (\frac{1}{w_i} - 1) - m_i w_i + w_i y_i) x_i\} = -\sum_{i=1}^N (y_i - m_i w_i) x_i$$

In matrix form:  $\nabla l(\beta) = -X'(y - mw) = X'(mw - y)$

## Part B

The following is my gradient descent code and results. A few notes regarding methodology:

- (a) Step size is fixed at  $a = .01$  in this code.
- (b) To handle probabilities close to 1 and 0, .01 is added to each log term in the loglikelihood.
- (c) Intercept was handled by adding a column of 1's to the X matrix.
- (d) Convergence was determined using  $\|\nabla l(\beta)\| < 1 * 10^{-2}$ .

R output results were as follows.

```
[1] "Algorithm has converged."
[1] 20311
>
> #Post-processing steps.
5 > beta_gd <- betas[[iter]] #Save and output estimated beta values.
> beta_gd
```

	V3	V4	V5	V6	V7	V8
V9	0.43486932	-5.03700599	1.65564614	-3.63350426	13.54124216	1.05599002
	0.69689744					0.02689289

```

10 > beta_gd
      V10      V11      V12
2.61356777 0.44381978 -0.48663758
> beta #Output glm beta values for comparison.
```

	X	XV3	XV4	XV5	XV6	XV7	XV8
XV9	0.48701675	-7.22185053	1.65475615	-1.73763027	14.00484560	1.07495329	-0.07723455
	0.67512313						

```

15 > beta_gd
      XV10      XV11      XV12
2.59287426 0.44625631 -0.48248420
```




Figure 4: Log-likelihood function

The R code is as follows.

```

### SDS 385 – Exercises 01 – Part B – Problem B
#This code implements gradient descent to estimate the
#beta coefficients for binomial logistic regression.

5 #Jennifer Starling
  #26 August 2016

library(Matrix)
rm(list=ls())

10 #PART B:

#Read in code.
wdbc = read.csv('/Users/jennstarling/UTAustin/2016_Fall_SDS 385_Stats Models for
  Big Data/Course Data/wdbc.csv', header=FALSE)
15 y = wdbc[,2]

#Convert y values to 1/0's.
Y = rep(0,length(y)); Y[y=='M']=1
X = as.matrix(wdbc[, -c(1,2)])

20 #Select features to keep, and scale features.
scrub = which(1:ncol(X) %% 3 == 0)
scrub = 11:30
X = X[, -scrub]
25 X <- scale(X) #Normalize design matrix features.
X = cbind(rep(1,nrow(X)),X)

```



```

#Set up vector of sample sizes. (All 1 for wdbc data.)
m <- rep(1,nrow(X))

30 #-----
#Binomial Negative Loglikelihood function.
#Inputs: Design matrix X, vector of 1/0 vals Y,
# coefficient matrix beta, sample size vector m.
35 #Output: Returns value of negative log-likelihood
# function for binomial logistic regression.
logl <- function(X,Y,beta,m){
  w <- 1 / (1 + exp(-X %*% beta)) #Calculate probabilities vector w_i.
  logl <- - sum(Y*log(w+1E-4) + (m-Y)*log(1-w+1E-4)) #Calculate log-likelihood.
40 #Adding constant to resolve issues with probabilities near 0 or 1.
  return(logl)
}

#-----
45 #Function for calculating Euclidean norm of a vector.
norm_vec <- function(x) sqrt(sum(x^2))

#-----
#Gradient Function:
50 #Inputs: Design matrix X, vector of 1/0 vals Y,
# coefficient matrix beta, sample size vector m.
#Output: Returns value of gradient function for binomial
# logistic regression.

55 gradient <- function(X,Y,beta,m){
  w <- 1 / (1 + exp(-X %*% beta)) #Calculate probabilities vector w_i.

  gradient <- array(NA,dim=length(beta)) #Initialize the gradient.
  gradient <- -apply(X*as.numeric(Y-m*w),2,sum) #Calculate the gradient.
60
  return(gradient)
}

#-----
65 #Gradient Descent Algorithm:
#Inputs:
# X: n x p design matrix.
# Y: response vector length n.
# m: vector length n.
70 # conv: Tolerance level for determining convergence, (length of gradient) < conv
# a: Step size.

#Outputs:
# beta_hat: A vector of estimated beta coefficients.
75 # iter: The number of iterations until convergence.
# converged: 1/0, depending on whether algorithm converged.
# loglik: Log-likelihood function.

```

```

gradient_descent <- function(X,Y,m,maxiter=50000,conv=1E-10,a=.01){
80
  #1. Initialize values.
  loglik <- rep(0,maxiter)      #Initialize vector to hold loglikelihood function.

  #Initialize matrix to hold gradients for each iteration.
85  grad <- matrix(0,nrow=maxiter,ncol=ncol(X))

  #Initialize matrix to hold beta vector for each iteration.
  betas <- matrix(0,nrow=maxiter+1,ncol=ncol(X))

90  converged <- 0      #Indicator for whether convergence met.

  #Initialize values for first iteration.
  betas[1,] <- rep(0,ncol(X)) #Initialize beta vector to 0 to start.
  loglik[1] <- logl(X,Y,betas[1,],m)
95  grad[1,] <- gradient(X,Y,betas[1,],m)

  #2. Perform gradient descent.
  for (i in 2:maxiter){

100    #Set new beta equal to beta - a*gradient(beta).
    betas[i,] <- betas[i-1,] - a * grad[i-1,]

    #Calculate loglikelihood for each iteration.
    loglik[i] <- logl(X,Y,betas[i,],m)

105    #Calculate gradient for beta.
    grad[i,] <- gradient(X,Y,betas[i,],m)

    #Check if convergence met: If yes, exit loop.
110    if (abs(loglik[i]-loglik[i-1])/abs(loglik[i-1]+1E-3) < conv){
      converged=1;
      break;
    }

115  } #End gradient descent iterations.

  return(list(beta_hat=betas[i,], iter=i, converged=converged, loglik=loglik[1:i
    ]))
}

120 #-----
#Run gradient descent and view results.

#1. Fit glm model for comparison. (No intercept: already added to X.)
glm1 = glm(y~X-1, family='binomial') #Fits model, obtains beta values.
125 beta <- glm1$coefficients

#2. Call gradient descent function to estimate.
beta_hat <- gradient_descent(X,Y,m,maxiter=100000,conv=1E-10,a=.01)

130 #3. Eyeball values for accuracy & display convergence.

```

```
beta          #Glm estimated beta values.
beta_hat$beta_hat  #Gradient descent estimated beta values.

print(c("Algorithm converged? ",beta_hat$converged, " (1=converged, 0=did not
      converge)"))
135 print(beta_hat$iter)

#4. Plot log-likelihood function for convergence.
plot(1:length(beta_hat$loglik),beta_hat$loglik,type='l',xlab='iterations',col='
      blue',log='xy')

140 #Save plot.
jpeg(file='/Users/jennstarling/UTAustin/2016_Fall_SDS 385_Stats Models for Big
      Data/385_Exercise_R_Code/R_Output/Ex01_B_loglik.jpeg')
plot(1:length(beta_hat$loglik),beta_hat$loglik,type='l',xlab='iterations',col='
      blue')
dev.off()
```

## Part C

### Hessian:

First, find the Hessian of  $l(\beta)$ , as this will be a key part of the Taylor series expansion.

$$\nabla^2 l(\beta) = \frac{\delta}{\delta \beta} [-\sum_{i=1}^N (y_i x_i - m_i w_i x_i)] = \sum_{i=1}^N m_i x_i \left( \frac{\delta w_i}{\delta \beta} \right).$$

From previous parts,  $\frac{\delta w_i}{\delta \beta} = x_i w_i^2 \exp(-x_i' \beta)$ , and  $\exp(-x_i' \beta) = (\frac{1}{w_i} - 1)$ , so plug in to get  $\frac{\delta w_i}{\delta \beta} = x_i w_i^2 (\frac{1}{w_i} - 1) = x_i w_i (1 - w_i)$ .

$$\text{Then } \nabla^2 l(\beta) = \frac{\delta^2 l}{\delta \beta \delta \beta'} = \sum_{i=1}^N m_i x_i x_i' w_i (1 - w_i).$$

In matrix form,  $\nabla^2 l(\beta) = X' A X$ , with  $A$  = diagonal matrix of  $m_i w_i (1 - w_i)$  elements.

### Taylor Series Second-Order Expansion:

General multivariate 2nd order Taylor series form:  $q(x; a) = f(a) + g(a)'(x - a) + \frac{1}{2}(x - a)' H(a)(x - a)$  where  $g(a)$  indicates the gradient evaluated at  $a$ , and  $H(a)$  indicates the Hessian evaluated at  $a$ .

$$q(\beta; \beta_0) = l(\beta_0) + g'(a)(\beta - \beta_0) + \frac{1}{2}(\beta - \beta_0)' H(\beta - \beta_0)$$

Plug in  $g(a) = X'(mw - y)$  and  $H = X' A X$ :  $q(\beta; \beta_0) = l(\beta_0) + [X'(mw - y)]'(\beta - \beta_0) + \frac{1}{2}(\beta - \beta_0)' X' A X (\beta - \beta_0)$

Distribute the transpose in the middle term:  $q(\beta; \beta_0) = l(\beta_0) + (Y - mw)' X (\beta - \beta_0) + \frac{1}{2}(\beta - \beta_0)' X' A X (\beta - \beta_0)$

Expand terms:  $q(\beta; \beta_0) = l(\beta_0) + (Y - mw)' X \beta + (Y - mw)' X \beta_0 + \frac{1}{2} \beta' X' A X \beta - 2(\frac{1}{2}) \beta_0' X' A X \beta + \frac{1}{2} \beta_0' X' A X \beta_0$

The following terms are constant, so let  $C = l(\beta_0) + (Y - mw)' X \beta_0 + \frac{1}{2} \beta_0' X' A X \beta_0$ .

Then rewrite as:  $q(\beta; \beta_0) = C + (Y - mw)' X \beta - \beta_0' X' A X \beta + \frac{1}{2} \beta' X' A X \beta$

Group first-order terms, treating  $X\beta$  as the variable in the quadratic form:

$$q(\beta; \beta_0) = C + [(Y - mw)' - A X \beta_0]' X \beta + \frac{1}{2} (X \beta)' A X \beta$$

(a) In second term, brought  $\beta_0' X' A$  into transpose, so  $(\beta_0' X' A)' = A' X \beta_0 = A X / \beta_0$  since  $A' = A$ .

(b) In third term, rearranged so  $\beta' X' = (X \beta)'$

Now complete the square, using trick:

$$a + b'X + X'CX = \frac{1}{2}(X - m)'M(X - m) \text{ with } M = C, m = -C^{-1}, v = a - \frac{1}{2}b'C^{-1}b$$

For our equation:

(a) In the expanded form, we have  $a = \text{constant}$ ,  $b = [(y - mw)' - A X \beta_0]$ ,  $c = A$ .

(b) Then our  $M = A$ ,  $v = \text{constant}$  that does not depend on  $\beta$ ,  $m = -A^{-1}[(y - mw)' - A X \beta_0] = [A^{-1}(y - mw) + X \beta_0] = Z$

Therefore,  $q(\beta; \beta_0) = \frac{1}{2}(X \beta - Z)' A (X \beta - Z) + C = \frac{1}{2}(Z - X \beta)' A (Z - X \beta) + C$ , where:

(a)  $A$  = diagonal matrix, with diagonal elements  $m_i w_i (1 - w_i)$

(b)  $Z = [A^{-1}(y - mw) + X \beta_0]$

(c)  $C$  = a constant that does not depend on  $\beta_0$

## Part D

My implementation of the Newton method converged in 10 iterations. Results are as follows.

```

> beta                                     #Output glm beta values for comparison.
      X      XV3      XV4      XV5      XV6      XV7      XV8
XV9
 0.48701675 -7.22185053  1.65475615 -1.73763027 14.00484560  1.07495329 -0.07723455
0.67512313
      XV10      XV11      XV12
2.59287426  0.44625631 -0.48248420
> beta_newt
[1]  0.48701675 -7.22185053  1.65475615 -1.73763027 14.00484560  1.07495329 -0.07723455
0.67512313
[9]  2.59287426  0.44625631 -0.48248420

```




Figure 5: Log-likelihood for Newton algorithm

The R code for implementing the Newton algorithm is as follows.

```

### SDS 385 - Exercises 01 - Part B - Problem D
#This code implements Newton's Method to estimate the
#beta coefficients for binomial logistic regression.

#Jennifer Starling
#26 August 2016

rm(list=ls())
library(Matrix)

#PART C:

```

```

#Read in code.
wdbc = read.csv('/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Stats Models for
  Big Data/Course Data/wdbc.csv', header=FALSE)
15 y = wdbc[,2]

#Convert y values to 1/0's.
Y = rep(0,length(y)); Y[y=='M']=1
X = as.matrix(wdbc[, -c(1,2)])
20

#Select features to keep, and scale features.
scrub = which(1:ncol(X) %% 3 == 0)
scrub = 11:30
X = X[, -scrub]
25 X <- scale(X) #Normalize design matrix features.
X = cbind(rep(1,nrow(X)),X)

#Set up vector of sample sizes. (All 1 for wdbc data.)
m <- rep(1,nrow(X))
30

#-----
#Binomial Negative Loglikelihood function.
#Inputs: Design matrix X, vector of 1/0 vals Y,
# coefficient matrix beta, sample size vector m.
35 #Output: Returns value of negative log-likelihood
# function for binomial logistic regression.
logl <- function(X,Y,beta,m){
  w <- 1 / (1 + exp(-X %*% beta)) #Calculate probabilities vector w_i.
  logl <- - sum(Y*log(w+1E-4) + (m-Y)*log(1-w+1E-4)) #Calculate log-likelihood.
40 #Adding constant to resolve issues with probabilities near 0 or 1.
  return(logl)
}

#-----
45 #Function for calculating Euclidean norm of a vector.
norm_vec <- function(x) sqrt(sum(x^2))

#-----
#Gradient Function:
50 #Inputs: Design matrix X, vector of 1/0 vals Y,
# coefficient matrix beta, sample size vector m.
#Output: Returns value of gradient function for binomial
# logistic regression.
55 gradient <- function(X,Y,beta,m){
  w <- 1 / (1 + exp(-X %*% beta)) #SCalculate probabilities vector w_i.

  gradient <- array(NA,dim=length(beta)) #Initialize the gradient.
  gradient <- -apply(X*as.numeric(Y-m*w),2,sum) #Calculate the gradient.
60
  return(gradient)
}

#-----

```

```

65 #Gradient Function:

hessian <- function(X,Y,beta,m){
  w <- 1 / (1 + exp(-X %*% beta)) #Calculate probabilities vector w_i.

70  #Create diag matrix of weights with ith element equal to m_i*w_i*(1-w_i)
  A <- Diagonal(length(m),m*w*(1-w))

  #Calculate Hessian as X'AX.
  H <- t(X) %*% A %*% X
75  return(H)
}

#-----
#QR Solver Function:
80 qr_decomp <- function(A,b){
  #Solves linear system Ax=b.

  #Obtain QR decomposition of matrix A. Extract components.
  QR <- qr(A)
85  Q <- qr.Q(QR)
  R <- qr.R(QR)

  #Backsolve for x.
  x <- qr.solve(A,b)
90  return(x)
}

#-----

cholesky_method <- function(X,W,y){
95  #Solves linear system Ax=b.
  #Since we have (X'WX)B=X'Wy, B (beta) acts as x, with A and b as follows.

  #Finding B (beta_hat) in equation
  A = (t(X) * diag(W)) %*% X #Efficient way of A = t(X) %*% W %*% X as W diag.
100  #Avoids mult by 0's.
  b = (t(X) * diag(W)) %*% y #b'Wy
  R <- chol(A) #Find right/upper cholesky decomposition of A.

  #Now we have R'R=A.
105  #1. Solve R'z=b for z. This is z = inv(R')b.
  z = solve(t(R)) %*% b

  #2. Solve Rx=z for x. This is x = inv(R)z. (x = beta_hat)
110  B_hat_chol <- solve(R) %*% z

  return(B_hat_chol)
}

#-----
115 #Newton's Method algorithm:

```

```
#1. Fit glm model for comparison. (No intercept: already added to X.)
glm1 = glm(y~X-1, family='binomial') #Fits model, obtains beta values.
120 beta <- glm1$coefficients

loglik <- 0           #Initialize vector to hold loglikelihood function.
grad <- list()        #Initialize list to hold gradients for each iteration.
hess <- list()        #Initialize list to hold hessians for each iteration.
125 maxiter <- 100000  #Specify max iterations allowed.
betas <- list()       #Initialize list to hold beta vector for each iteration.

conv <- 1E-6          #Set convergence level.

130 #Initialize first iteration of values.
```



**Part E**

Newton is a second-order optimization method, while Gradient Descent is a first-order. So Newton uses the second derivatives in determining the direction to take each step. This means that Newton takes fewer iterations to find the local min of the cost function than gradient descent.

However, iteration-to-iteration, Newton is a more expensive function to calculate. It requires evaluation of the Hessian matrix, and solving the linear system  $Hessian * dir = Gradient$  for the dir (direction) gradient, which is used to update the betas for the next step:

$$\beta^{(i+1)} = \beta^{(i)} - H^{-1}(\beta^{(n)}) \nabla(\beta^{(n)})$$

Instead of inverting the Hessian directly, we can solve the following equation with a QR solver (or other matrix decomposition method).

$$H \cdot dir = \nabla, \text{ which yields } dir = H^{-1} \nabla, \text{ and so } \beta^{(i+1)} = \beta^{(i)} - dir.$$

This does improve the speed of the Hessian inversion, but is still an added cost compared to the calculation of each Gradient Descent iteration.

The addition of the Hessian matrix may also be problematic for the Newton method; if the Hessian is singular, the method does not work.