

**SDS 385 Ex 04:**  
**Lazy Updates for L2 Regularization in Adagrad**

September 30, 2016

Jennifer Starling

## Problem:

For sparse matrices, any of the coefficients  $\beta_j$  are not necessarily updated at every iteration. Only active (non-zero)  $\beta_j$ 's are updated at a given iteration. However, even if a  $\beta_j$  is not updated, the L2 regularization penalty still accumulates over each iteration.

The next time  $\beta_j$  is updated, before the update occurs, the accumulated L2 regularization penalty must be added to the  $\beta_j$ .

Example:

Iteration	Update
10	$\beta_j$ updated
11	No update to $\beta_j$
12	No update to $\beta_j$
13	No update to $\beta_j$
14	$\beta_j$ updated

For iteration 14, two things must happen:

- (1) Add the accumulated regularization penalty for iterations 11, 12 and 13.
- (2) Then update  $\beta_j$  like normal.

## L2 Regularization Penalty for A Single Iteration

For a single iteration,  $\beta_j$  is updated as follows.

$$\beta_j^{(i+1)} = \beta_j^{(i)} - \lambda \gamma^{(i)} \beta_j^{(i)} \text{ where:}$$

$\beta_j^{(i)}$  is the  $\beta_j$  from the previous iteration.

$\lambda$  is the L2 regularization  $\lambda$  penalty coefficient.

$\gamma^{(i)}$  is the  $(step * adj.grad_j)$ , the previous Adagrad scaled step size.

Recall,  $adj.grad_j = grad_j * invSqrt(hist.grad_j + \epsilon)$ , where  $hist.grad_j = \text{running sum of } (grad_j)^2$ .

## L2 Regularization Penalty for Consecutive Iterations

Let  $\beta_j^{(1)}$  be a normal update. Let  $\gamma^{(1)}$  be the  $adj.grad_j$  corresponding to this update.

A key feature of Adagrad is that  $adj.grad_j$  only changes when  $\beta_j$  is updated.

For subsequent iterations, if  $\beta_j$  were being updated, the recursion that occurs is as follows.

$$\begin{aligned}
 \beta_j^{(2)} &= \beta_j^{(1)} - \lambda \gamma^{(1)} \beta_j^{(1)} = (1 - \lambda \gamma^{(1)}) \beta_j^{(1)} \\
 \beta_j^{(3)} &= \beta_j^{(2)} - \lambda \gamma^{(1)} \beta_j^{(2)} = (1 - \lambda \gamma^{(1)}) \beta_j^{(1)} - \lambda \gamma^{(1)} (1 - \lambda \gamma^{(1)}) \beta_j^{(1)} = (1 - \lambda \gamma^{(1)})^2 \beta_j^{(1)} \\
 &\vdots \\
 \beta_j^{(i+1)} &= (1 - \lambda \gamma^{(1)})^i \beta_j^{(1)}
 \end{aligned}$$

## Summing Consecutive Iteration Penalties

The above recursion yields a  $\beta_j$  update for each iteration after the first one. Must sum these terms for the iterations where an update did not occur before proceeding with the update as usual.

- Let  $i = 0$  the last update;  $\beta_j^{(0)}$  is the last regularly updated  $\beta_j$ , with corresponding  $\gamma^{(0)}$ .
- Let *curr.iter* represent the current iteration where  $\beta_j$  is now being updated again.
- Let *last.update* =  $i = 0$ , and call *skip* = *curr.iter* - *last.update*.

Sum the penalty terms since the last known update:

$$\beta_j^\Delta = \beta_j^{(1)} + \beta_j^{(2)} + \beta_j^{(3)} + \dots$$

$$\beta_j^\Delta = \sum_{i=0}^{(skip-1)} (1 - \lambda\gamma^{(0)})^{i+1} \beta_j^{(0)}$$

The summation notation is effectively beginning with  $\beta_j^{(1)}$  by raising to the power  $(i + 1)$ .

The summation stops at *curr.iter* - *last.update* - 1 = *skip* - 1 because we want to accumulate penalties up to and including the current  $\beta_j$  update.

Example:

- If the last update was iteration 10, and the current update is iteration 15, want to add penalties for 11, 12, 13 and 14.
  - The penalty for iteration 15 will occur as part of the regular update of  $\beta_j$  in iteration 15.
  - *curr.iter* - *last.update* = 15 - 10 = 4 here, so *skip* = 4.
  - In this example,
    - iteration 10 is  $i=0$ , so we raise to the  $(i+1)=1$  power for  $(i=0)$ , adding the 11th iteration penalty.
    - iteration 11 is  $i=1$ , so we raise to the  $(i+1)=2$  power for  $(i=1)$ , adding the 12th iteration penalty.
    - iteration 12 is  $i=2$ , so we raise to the  $(i+1)=3$  power for  $(i=2)$ , adding the 13th iteration penalty.
    - iteration 13 is  $i=3$ , so we raise to the  $(i+1)=4$  power for  $(i=3)$ , adding the 14th iteration penalty.
    - Now we are at  $i=4$ , but we have already added the 14th iteration penalty.
- This is why we are stopping at *skip* - 1 instead of stopping at *skip*.

Note that the summation just reuses  $\beta_j^{(0)}$  and  $\gamma^{(0)}$  from the last update; it is not required to know any interim  $\beta_j$  updates, thanks to the recursion.

The summation  $\beta_j^\Delta = \sum_{i=0}^{(skip-1)} (1 - \lambda\gamma^{(0)})^{i+1} \beta_j^{(0)}$  has the form  $\beta_j^{(0)} * \sum_{i=0}^n r^k = \beta_j^{(0)} \frac{1-r^{n+1}}{1-r}$  where

- $r = 1 - \lambda\gamma^{(0)}$
- $(n + 1) = (skip - 1) + 1 = skip$

## Result

The cumulative L2 penalty term that must be added prior to the current update of  $\beta_j$  is:

$$\beta_j^\Delta = \beta_j^{(0)} \left( \frac{1 - (1 - \lambda\gamma^{(0)})^{skip}}{\lambda\gamma^{(0)}} \right) \text{ with:}$$

- $\beta_j^{(0)}$  = last updated  $\beta_j$
- $\gamma^{(0)}$  = scaled adagrad step corresponding to last updated  $\beta_j$ .
- $step = current.iteration.number - last.updated.iteration.number$  for the  $\beta_j$  being updated.

## Note regarding Minimization vs Maximization

The above result is for maximizing the gradient of the likelihood. If you are minimizing the gradient of the negative log-likelihood (as I am doing in my C++ code), the  $\beta_j^\Delta$  term changes as follows.

For minimizing the gradient of the negative log-likelihood, updates to the gradient, including the penalty, are:

$$grad_j^{(i+1)} = (m_i * w_i - Y_i)X_{ij} + 2\lambda\beta_j^{(i)}$$

In this case, the penalty term is being added instead of subtracted.

Therefore, the summation  $\beta_j^\Delta = \sum_{i=0}^{(skip-1)} (1 + \lambda\gamma^{(0)})^{i+1} \beta_j^{(0)}$  has the form  $\beta_j^{(0)} * \sum_{i=0}^n r^k = \beta_j^{(0)} \frac{1-r^{n+1}}{1-r}$  where

- $r = 1 + \lambda\gamma^{(0)}$
- $(n+1) = (skip-1) + 1 = skip$

The summation therefore simplifies to:

$$\beta_j^\Delta = \beta_j^{(0)} \left( \frac{1 - (1 + \lambda\gamma^{(0)})^{skip}}{1 - \lambda\gamma^{(0)}} \right) \text{ with:}$$

- $\beta_j^{(0)}$  = last updated  $\beta_j$
- $\gamma^{(0)}$  = scaled adagrad step corresponding to last updated  $\beta_j$ .
- $step = current.iteration.number - last.updated.iteration.number$  for the  $\beta_j$  being updated.

Pseudo-code to illustrate this change is as follows, to provide clarity on correct signs of operations.

```
//Step 1: Apply accumulated l2 penalty to beta_hat_j.
//Skip = Number of iters since last update. (Skip=1 means updated last iter.)
skip = iter - last_updated(j);
5
//Calculate accum penalty. Based on recursion defined in my notes.
gam = step*adj_grad(j);
accum_l2_penalty = beta_hat(j) * ((1-pow(1+lambda*gam,skip))/(1-lambda*gam));
10
//Add accum l2 penalty to beta_hat_j before doing current iteration update.
beta_hat(j) -= accum_l2_penalty;

//Step 2: Continue with updates for jth row in ith col.
15
//Calculate l2 norm penalty.
double l2penalty = 2*lambda*beta_hat(j);

//Update the jth gradient term. Note: it.value() looks up Xji for nonzero entries.
grad_j = (mi*wi-Yi) * it.value() + l2penalty;
20
//Update the jth hist_grad term for Adagrad.
hist_grad(j) += grad_j * grad_j;

//Calculate the jth adj_grad term for Adagrad.
25
adj_grad(j) = grad_j * invSqrt(hist_grad(j) + epsilon);

//Calculate the updated jth beta_hat term.
beta_hat(j) -= step*adj_grad(j);
```