

**SDS 385 Ex 05:**  
**Sparsity in Covariates (LASSO)**

October 3, 2016

**Jennifer Starling**

## Problem 1 - Penalized Regression & Soft Thresholding

### Part A

#### Proof 1

Show the quadratic term in the objective is the negative loglikelihood of a Gaussian with mean  $\theta$ , var = 1.

The *Gaussian*( $\theta, 1$ ) distribution has pdf  $f(y|\theta, 1) = (2\pi\sigma^2)^{-1/2} \exp[-\frac{(y-\theta)^2}{2\sigma^2}] = (2\pi)^{-1/2} \exp[-\frac{(y-\theta)^2}{2}]$

Then the likelihood function is  $L(\theta|y, 1) = (2\pi)^{-1/2} \exp[-\frac{(y-\theta)^2}{2}]$  (looking at a single y)

Then the log-likelihood function is  $\log L = \log(L(\theta|y, 1)) = \frac{-1}{2} \log(2\pi) - \frac{1}{2}(y - \theta)^2$ , and the first term is not dependent on  $\theta$ , so we drop it out.

Then the negative log-likelihood function is  $-\frac{1}{2}(y - \theta)^2$ .

Therefore, the quadratic term in the objective is the negative log-likelihood of a Gaussian with mean  $\theta$  and variance 1. ■

#### Proof 2

Prove that  $S_\lambda(y) = \text{sign}(y)(|y| - \lambda)_+$  where  $a_+ = \max(a, 0)$ , the positive part of  $a$ .

Take the derivate of  $S_\lambda(y)$  to obtain:

$$\frac{\partial S_\lambda(y)}{\partial \theta} = \frac{\partial}{\partial \theta} \left( \frac{1}{2}(y - \theta)^2 + \lambda|\theta| \right) = -(y - \theta) + \lambda \frac{|\theta|}{\theta} = -(y - \theta) + \lambda * \text{sign}(\theta)$$

Break the problem into three cases:

- (1)  $\theta > 0$
- (2)  $\theta < 0$
- (3)  $\theta = 0$

Case 1: If  $\theta > 0$ , the objective function set equal to zero can be rewritten as  $-(y - \theta) + \lambda = 0 \rightarrow \theta = y - \lambda$

Constraint:  $\theta > 0$ , and  $\theta = y - \lambda$ , so  $y - \theta > 0 \rightarrow y > \lambda$

Case 2: If  $\theta < 0$ , the objective function set equal to zero can be rewritten as  $-(y - \theta) - \lambda = 0 \rightarrow \theta = y + \lambda$

Constraint:  $\theta < 0$ , and  $\theta = y + \lambda$ , so  $y + \lambda < 0 \rightarrow y < -\lambda$

Case 3: The above cases cover  $y > \lambda$  and  $y < -\lambda$ , leaving  $|y| < \lambda$  as the  $\theta = 0$  constraint.

These three cases can be summarized in a single function:

$$S_\lambda(y) = \text{sign}(y)(|y| - \lambda)_+ \quad \blacksquare$$

**Part B**1 through 3

The plotted grid of  $\hat{\theta}(y_i)$  versus  $\theta_i$  across varying  $\lambda$  values shows how certain  $\theta_i$ 's are selected (i.e. set to zero), while the non-zero  $\theta_i$ 's are shrunk towards zero.

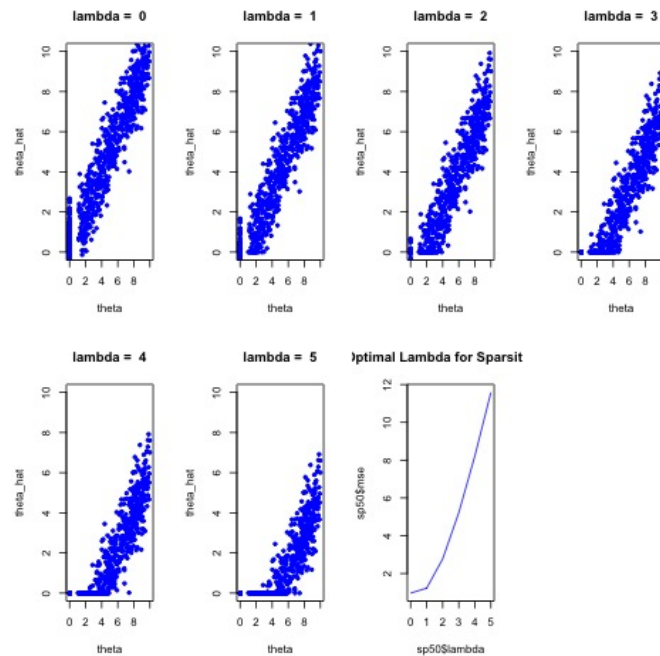
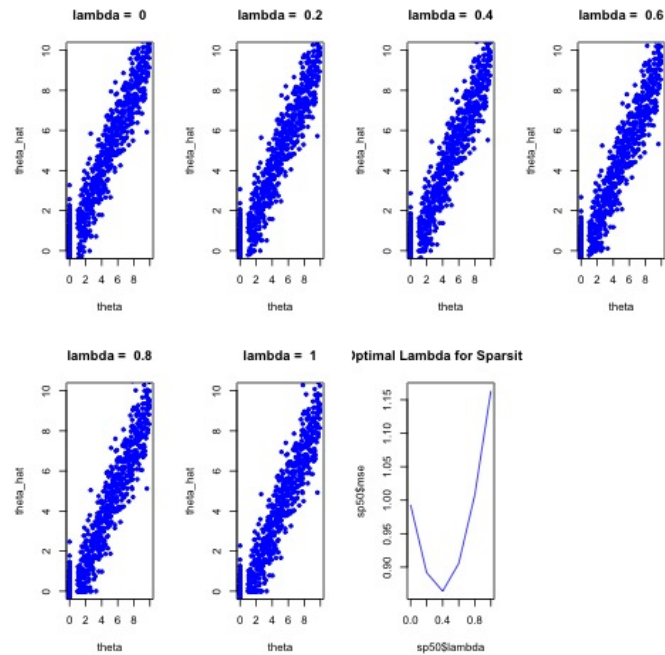
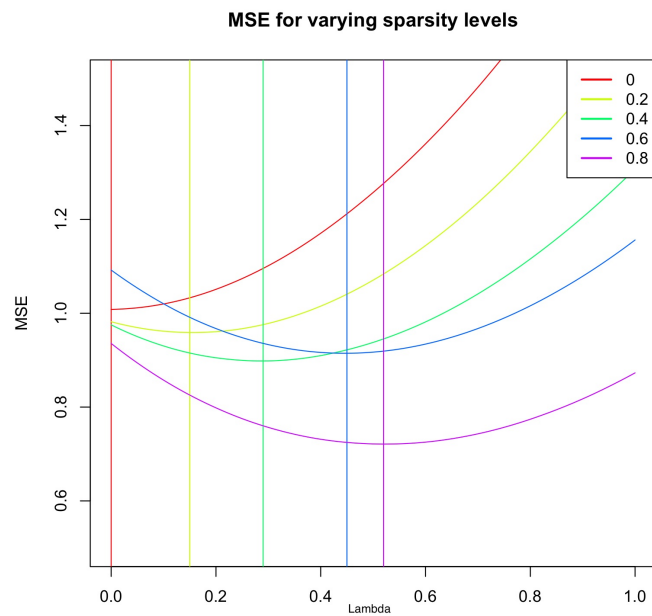


Figure 1:  $\hat{\theta}(y_i)$  vs  $\theta_i$  for varying  $\lambda$  at 50% sparsity

Below is a similar figure, for a smaller group of lambdas (0.0, 0.2, 0.4, 0.6, 0.8, 1.0), which better illustrates finding an optimal lambda value. This plot is also done using 50% sparsity.

Figure 2:  $\hat{\theta}(y_i)$  vs  $\theta_i$  for varying  $\lambda$  at 50% sparsity

4 The plot of the MSE of the estimate as a function of  $\theta$  is as follows. Horizontal lines represent the location of the minimum MSE for each sparsity level. Notice that increasing sparsity has the effect of increasing the optimal lambda value.

Figure 3: MSE as a function of  $\lambda$  for varying levels of sparsity

## Problem 2 - The LASSO

### Part A

The plot of the solution path of  $\hat{\beta}_\lambda$  as a function of  $\lambda$  is as follows. Note that the glmnet plot function plots based on  $\log(\lambda)$  instead of  $\lambda$  so that the  $\lambda$  values are evenly spaced. The values across the top of the plot represent degrees of freedom, ie non-zero coefficients, at each  $\lambda$  value.




Figure 4: solution path of  $\hat{\beta}_\lambda$  as a function of  $\lambda$

The optimal lambda in this case was zero, which makes sense because the model is being optimally fit to the  $\hat{\beta}$ 's.

See section C for the plot of the MSE, MOOSE (CV) and Cp errors together.

## **Part B**

The optimal lambda selected by my custom cross-validation function was 0.03493409. (This matched the lambda using the glmnet built-in cv functionality.)

See section C for the plot of the MSE, MOOSE (CV) and Cp errors together.

## Part C

The optimal lambda selected by Mallows's Cp was 0.03493409. This matched the optimal lambda selected by cross-validation in part B.

The plots of the MSE, CV (MOOSE) and Cp errors by lambda is below. The second plot is using  $\log(\text{lambda})$  on the x-axis. The optimal lambda occurs at the same value for CV and Cp.


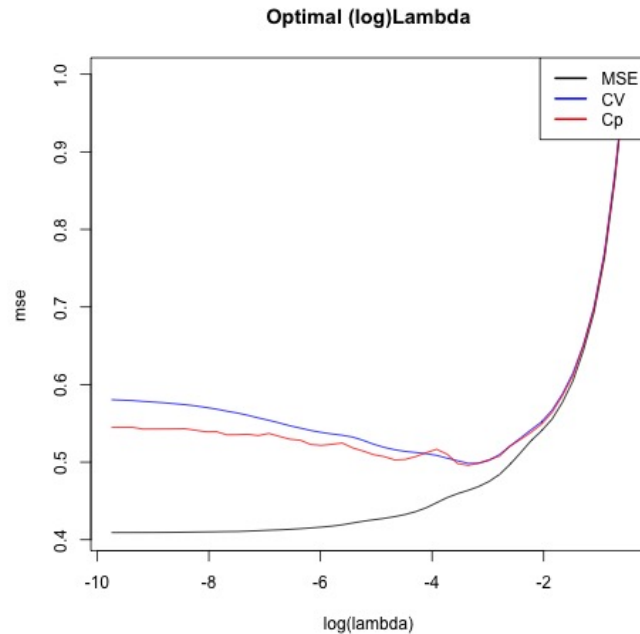


Figure 5: Various errors by  $\lambda$

Figure 6: Various errors by  $\log(\lambda)$ 

## Appendix: R Code

```

#Big Data Exercise 5
#Jennifer Starling
#27 Sept 2016

5 rm(list=ls())

#####
### PENALIZED LHOOD & SOFT THRESHOLDING ###
#####

10 #
#Functionalize the toy example, so that is easy to call for various sparsity
#levels.
toy_example <- function(n,sd,sparsity,lambda){
#Function inputs:
15 #n = sample size
#sd = n-length vector of standard deviations.
#sparsity = percent of theta value sparse; 0 to 1.
#lambda = vector of lambda values to test.
#Function output:
20 #mse = Vector of MSE values for each theta.
#theta = 'True values' of generated thetas.
#theta_hat = Matrix of estimated theta values. Each col = a different lambda.

#Generate "true" theta values; different for each obs.
25 theta = sample(seq(1,10,by=.01),n,replace=T)
theta[sample(1:n,sparsity*n,replace=T)]=0

#Simulate n y-values, using yi ~ N(theta_i,sd_i)

```



```

y = rep(0,n)
30 for (i in 1:n){
    y[i] = rnorm(1,theta[i],sd[i])
}

#Initialize theta_hat matrix.
35 theta_hat <- matrix(0,nrow=n,ncol=length(lambda))
colnames(theta_hat) = paste('lambda=',lambda)

#Initialize vector to hold MSE for each lambda.
mse <- rep(0,length(lambda))

40 #Calculate theta_hat values for each lambda.
for (j in 1:length(lambda)){
    #Calculate the Sy function; takes a few steps.
    Sy <- abs(y)-lambda[j]
    Sy[which(Sy<0)]=0 #Take only the positive part.
45 Sy <- sign(y)*Sy #Multiply by sign of y.

    #Assign Sy=theta_hat to its column.
    theta_hat[,j] = Sy

50 #Calculate mse for lambda_j.
    mse[j] <- (1/n) * sum((theta_hat[,j] - theta)^2)
}

55 return(list(theta=theta,lambda=lambda,theta_hat=theta_hat,mse=mse,sparsity=
    sparsity))
} #end function.

#


---


60 #EXAMPLE 1: With exaggerated (large) lambdas, to show drastic shrinking/sparsity
    in coeffs.

#Try out function with 50% sparsity, n=1000, sd=1.
n=1000
sp50 = toy_example(n=1000,sd <- rep(1,n), sparsity=.5,lambda=c(0,1,2,3,4,5))

65 #Plot results for varying lambda values.
par(mfrow=c(2,4))
for (j in 1:length(sp50$lambda)){
    plot(sp50$theta,sp50$theta_hat[,j],col='blue',pch=20,xlim=c(0,10),ylim=c(0,10)
    ,
70 main=paste('lambda = ',sp50$lambda[j]),xlab='theta',ylab='theta_hat')
}

#Plot MSE for this sparsity level.
plot(sp50$lambda,sp50$mse,type='l',col='blue',main=paste('Optimal Lambda for
    Sparsity ',sp50$sparsity))

75 #


---


#EXAMPLE 2: With more realistic lambdas, sparsity 50%

sp50 = toy_example(n=1000,sd <- rep(1,n), sparsity=.5,lambda=seq(0,1,by=.2))

```

```

80 #Plot results for varying lambda values.
par(mfrow=c(2,4))
for (j in 1:length(lambda)){
  plot(sp50$theta,sp50$theta_hat[,j],col='blue',pch=20,xlim=c(0,10),ylim=c(0,10)
85   main=paste('lambda = ',sp50$lambda[j]),xlab='theta',ylab='theta_hat')
}

#Plot MSE for this sparsity level.
plot(sp50$lambda,sp50$mse,type='l',col='blue',main=paste('Optimal Lambda for
  Sparsity ',sp50$sparsity))

90 #

#PROBLEM B-4: Plot MSE for several configurations of theta (sparsity levels)
#and observe how optimal lambda changes.

95 #Initialize a vector of sparsity levels.
sp_levels <- seq(0,.8,by=.2)
mse <- list()
lambda=seq(0,1,by=.01)

100 #Loop through sparsity levels.
for (s in 1:length(sp_levels)){

  #Run toy example for given sparsity level. Save mse.
  temp = toy_example(n=1000,sd <- rep(1,n), sparsity=sp_levels[s],lambda=lambda)
105   mse[[s]] = temp$mse
}

#Plot MSE for each sparsity level.
colors <- rainbow(length(sp_levels))
110 plot(lambda,mse[[1]],col = colors[1],type='l',xlim=c(0,1),ylim=c(.5,1.5),
  main='MSE for varying sparsity levels',xlab='lambda',ylab='MSE')
abline(v=lambda[which(mse[[1]]==min(mse[[1]]))],col=colors[1])

for (j in 2:length(sp_levels)){
115   lines(lambda,mse[[j]],col = colors[j],type='l')
  abline(v=lambda[which(mse[[j]]==min(mse[[j]]))],col=colors[j])
}
labels = paste(sp_levels)
legend('topright',legend=labels,lwd=2,col=colors, bty != "n", bg='white')

120 #####
###   THE LASSO   ###
#####

125 library(glmnet)

#Read in Diabetes.csv data.
X <- read.csv(file='/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Big_Data/
  Exercise_05_R_Code/DiabetesX.csv',header=T)
y <- read.csv(file='/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Big_Data/
  Exercise_05_R_Code/DiabetesY.csv',header=F)

130 #Scale X and y.
X = scale(X)

```

```

y = scale(y)

#-----
#Part A:

#Fit lasso model across a range of lambda values (which glmnet does automatically)
.

#Plot the solution path beta_hat_lambda as a function of lambda.
140 myLasso <- glmnet(X,y,family='gaussian',nlambda=50)

#Plot of beta_hat as function of lambda.
jpeg(file='/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Big_Data/Exercise05
      LaTeX_Files/LassoBetaPaths.jpg')
plot(myLasso,xvar="lambda", main='Lasso: Solution path of Beta_hat as function of
      lambda')
145 dev.off()

#Track in-sample MSE prediction error of the fit across the solution path:

lambda = myLasso$lambda
150 betas = myLasso$beta
n = nrow(X)

#Initialize vector to hold MSE for each beta.
MSE_betas = rep(0,length(lambda))

155 for (i in 1:length(lambda)){
      MSE_betas[i] = (1/n) * sum((y-X %*% betas[,i])^2)
}

160 jpeg(file='/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Big_Data/Exercise05
      LaTeX_Files/LassoMSE.jpg')
par(mfrow=c(1,2))
plot(log(lambda),MSE_betas,type='l')
plot(lambda,MSE_betas,type='l')
dev.off()

165 #-----
#Part B:

#Run cross-validation using vector of lambdas from model fit in part A.
170 lassoCV = myCV(X,y,lambda,cv_folds=10)

#Plot CV results to visualize optimal lambda.
plot(lassoCV$lambda,lassoCV$pred_err_by_lambda,type='l',col='blue')

175 #Output minimum lambda.
paste('Optimal lambda: ',
      lassoCV$lambda[which(lassoCV$pred_err_by_lambda==min(lassoCV$pred_err_by_
      lambda))])

#Test results against the built-in cross-val functionality in glmnet.
180 cvfit = cv.glmnet(X,y,lambda=lambda)
cvfit$lambda.min

#RESULTS:
#> paste('Optimal lambda: ',
185 #+ lassoCV$lambda[which(lassoCV$pred_err_by_lambda==min(lassoCV$pred_err_by_
      lambda))])

```

```

# [1] "Optimal lambda: 0.0349340915741447"
#> cvfit$lambda.min
# [1] 0.03493409
#
190 #My cross-validation function.
myCV <- function(X,y,lambda,cv_folds=10){
  #data = holds predictors and response.
  #lambda = vector of lambdas to include in the model.
195 #folds = number of cross-val folds.

  #Randomly shuffle data.
  data = cbind(y,X)

200 data = data[sample(nrow(data)),]

  #Create 'folds' number of equally sized folds.
  folds <- cut(seq(1,nrow(data)),breaks=cv_folds,labels=F)

205 #Initialize vector to hold prediction error for each cv fold iteration.
  pred_test_err = matrix(0,nrow=cv_folds,ncol=length(lambda))
  #pred_test_error <- rep(0,cv_folds)

  #Perform cross-validation.
210 for (i in 1:cv_folds){

    #Split up data using folds.
    testIndices <- which(folds==i,arr.ind=T)
    testData <- data[testIndices, ]
215 trainData <- data[-testIndices, ]

    #Fit glmnet lasso model for the data excluding the current fold.
    trainLasso = glmnet(x=trainData[,-1],y=trainData[,1],family='gaussian',
      nlambda=50,lambda=lambda)

220 #Predict values on test data.
    predLasso = predict(trainLasso,newx=testData[,-1],s=lambda)

    #Calculate and save prediction error.
    predErr = apply(predLasso,2,function(yhat) sum((yhat-testData[,1])^2))/
      nrow(testData)
225 pred_test_err[i,] = predErr
  }

  #Return average predicted test error for each k value.
  return(list(lambda=lambda,
230 pred_err_by_lambda=colMeans(pred_test_err),
    pred_err_var = apply(pred_test_err,2,var)))
}

#-----
235 #Part C: Compute and plot the Cp statistic (Mallow's Cp) as a function of lambda.

#Use the Part A glm lasso model, fit using whole data set for the same vector of
  50 lambdas.
#Also use the MSE calculated in part A.
lambda = myLasso$lambda
240 mse = MSE_betas
df = myLasso$df

```

```

n = nrow(X)

#Fit an OLS model to obtain estimate of sigma2.
245 mylm = lm(y~X-1)      #Fit model with no intercept
sigma2_hat = summary(mylm)$sigma^2

#Calculate Mallow's cp.
Cp = mse + 2 * (df/n) * sigma2_hat
250

#Plot Mallow's Cp as a function of lambda.
plot(lambda,Cp,type='l',main='Mallows Cp as a function of lambda',xlab='lambda',
      ylab='Cp')

#Output optimal lambda based on smallest Mallow's Cp.
255 paste('Optimal lambda based on Cp: ',lambda[which(Cp==min(Cp))])

#PLOTING:
jpeg(file='/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Big_Data/Exercise05
      LaTeX_Files/LassoOptimalLambda.jpg')
#Plot MSE, CV error and Cp on the same plot, to compare the optimal values of
      lambda they each yield.
260 plot(lambda,mse,type='l',col='black',main='Optimal Lambda')
lines(lambda,lassoCV$pred_err_by_lambda,col='blue')
lines(lambda,Cp,col='red')
labels = c('MSE','CV','Cp')
legend('topright',legend=labels,lwd=2,col=c('black','blue','red'), bty != "n", bg=
      'white')
265 dev.off()

jpeg(file='/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Big_Data/Exercise05
      LaTeX_Files/LassoOptimalLogLambda.jpg')
plot(log(lambda),mse,type='l',col='black',main='Optimal (log)Lambda')
lines(log(lambda),lassoCV$pred_err_by_lambda,col='blue')
270 lines(log(lambda),Cp,col='red')
labels = c('MSE','CV','Cp')
legend('topright',legend=labels,lwd=2,col=c('black','blue','red'), bty != "n", bg=
      'white')
dev.off()

```