

SDS 385 Ex 06:
The Proximal Gradient Method

October 8, 2016

Jennifer Starling

Proximal Operators

Part A

$$\hat{x} = \operatorname{argmin}_x \left\{ \tilde{l}(x; x_0) + \phi(x) \right\}$$

Plug into the Moreau envelope:

$$\begin{aligned} E_\gamma(\tilde{l}(x; x_0)) &= \operatorname{argmin}_x \left\{ l(x_0) + (z - x_0)^T \nabla l(x_0) + \frac{1}{2\gamma} (z - x_0)^T (z - x_0) + \phi(z) \right\} \leq \tilde{l}(x; x_0) \\ &= \operatorname{argmin}_x \left\{ l(x_0) + (z - x_0)^T \nabla l(x_0) + \frac{1}{2\gamma} (z^T z - 2z^T x_0 + x_0^T x_0) + \phi(z) \right\} \end{aligned}$$

Take derivative of the Moreau envelope and set equal to zero:

$$\nabla_z E_\gamma(\tilde{l}(x; x_0)) = \nabla l(x_0) + \frac{1}{2\gamma} (2z - 2x_0) + \phi(z) = 0$$

Solve for $\phi(z)$:

$$\phi(z) + \frac{z - x_0}{\gamma} = -\nabla l(x_0)$$

$$\phi(z) = x_0 - z - \gamma \nabla l(x_0)$$

To minimize, need $z = x_0 - \gamma \nabla l(x_0)$.

Therefore, the proximal operator solution has this form:

$$\hat{x} = \operatorname{prox}_\gamma(\phi(z)), \text{ where } z = x_0 - \gamma \nabla l(x_0).$$

Part B

1) Consider the Gaussian sampling model $(y|x) \sim N(Ax, \Omega^{-1})$.

The multivariate loglikelihood has form:

$$\log l(y) = \frac{1}{2} \log(|\Sigma|) + \frac{1}{2} (y - \mu)^T \Sigma^{-1} (y - \mu) + \frac{k}{2} \log(2\pi) \text{ for } N(\mu, \Sigma).$$

Plug into this form, and drop all terms which do not depend on x , to get:

$$l(x|y) = \frac{1}{2} (y - Ax)^T (y - Ax)$$

Then expand:

$$\begin{aligned} &= \frac{1}{2} (y^T - x^T A^T) \Omega (y - Ax) \\ &= \frac{1}{2} (y^T \Omega y - 2y^T \Omega Ax + x^T A^T \Omega Ax) \\ &= \frac{x^T A^T \Omega Ax}{2} - y^T \Omega Ax + \frac{y^T \Omega y}{2} \end{aligned}$$

Let:

- $P = A^T \Omega A$
- $q = -(y^T \Omega Ax)^T = -A^T \Omega y$
- $r = \frac{y^T \Omega y}{2}$

Then can write $l(x|y) = \frac{1}{2} x^T P x + q^T x + r$ ■

2) The proximal operator is derived as follows.

$$\begin{aligned} E_{\phi=1/\gamma}[l(x)] &= \operatorname{argmin}_z \left\{ l(z) + \frac{1}{1/\phi} \|z - x\|_2^2 \right\} \leq l(x) \\ &= \operatorname{argmin}_z \left\{ z^T P z + q^T z + r + \frac{1}{1/\phi} (z - x)^T (z - x) \right\} \leq l(x) \\ &= \operatorname{argmin}_z \left\{ z^T P z + q^T z + r + \frac{1}{1/\phi} (z^T z - z^T x + x^T x) \right\} \leq l(x) \end{aligned}$$

Then take derivative of the Moreau operator:

$$\nabla_x E_{\phi=1/\gamma}[l(x)] = \frac{1}{2} (2Pz) + q + \frac{1}{2\phi} (2x - 2z) = 0$$

$$Pz + q + \gamma(z - x) = 0, \text{ since } \phi = 1/\gamma$$

$$z(P + \gamma I) = \gamma x - q$$

$$z = (P - \gamma I)^{-1}(\gamma x - q)$$

Therefore, the proximal operator is $\operatorname{prox}_\phi = (P - \gamma I)^{-1}(\gamma x - q)$ ■

Part C

Let $\phi(x) = \tau\|x\|_1$. Express the proximal operator in terms of the soft-thresholding function from Exercise 5, $S_\lambda(y) = \text{sign}(y)(|y| - \lambda)_+$

- 1) First, the Moreau envelope is $E_\lambda(\phi(z)) = \underset{x}{\operatorname{argmin}} \left\{ \phi(z) + \frac{1}{2\gamma}\|z - x\|_2^2 \right\} \leq \phi(x)$
- 2) The Moreau envelope for a single element:

$$\tau|z| + \frac{1}{2\gamma}(z - x)^2$$

Case 1, if $z > 0$:

$$\frac{\delta}{\delta z} \left(\tau|z| + \frac{1}{2\gamma}(z - x)^2 \right) = \tau + \frac{1}{\gamma}(z - x) \rightarrow z = \left(\frac{x}{\gamma} - \tau \right) \gamma \rightarrow z = (x - \tau\gamma)$$

Case 2, if $z < 0$:

$$\frac{\delta}{\delta z} \left(\tau|z| + \frac{1}{2\gamma}(z - x)^2 \right) = -\tau + \frac{1}{\gamma}(z - x) \rightarrow z = \left(\frac{x}{\gamma} + \tau \right) \gamma \rightarrow z = (x + \tau\gamma)$$

Case 3, if $z = 0$:

The Moreau reduces to $\frac{1}{2\gamma}(-x)^2$, and the derivative with respect to z is zero.

These cases can be summarized by the soft thresholding function:

$$S_{\tau\gamma}(x) = \text{sign}(x)(|x| - \tau\gamma)_+$$

Therefore, the proximal operator $\text{prox}_{\tau\gamma}(x) = S_{\tau\gamma}(x) = \text{sign}(x)(|x| - \tau\gamma)_+$. ■

The Proximal Gradient Method

Part A

$$\hat{f}(x; x_0) = f(x_0) + (x - x_0)^T \nabla f(x_0)$$

1) Plug the approximation $\hat{f}(x; x_0)$ into the Moreau envelope:

$$\begin{aligned} E_\gamma(\hat{f}(x; x_0)) &= \operatorname{argmin}_x \left\{ \hat{f}(x; x_0) + \frac{1}{2\gamma} \|x - x_0\|_2^2 \right\} \leq f(x) \\ &= \operatorname{argmin}_x \left\{ f(x_0) + (x - x_0)^T \nabla f(x_0) + \frac{1}{2\gamma} (x - x_0)^T (x - x_0) \right\} \leq f(x) \\ &= \operatorname{argmin}_x \left\{ f(x_0) + (x - x_0)^T \nabla f(x_0) + \frac{1}{2\gamma} (x^T x - 2x^T x_0 + x_0^T x_0) \right\} \leq f(x) \end{aligned}$$

2) Take gradient of the Moreau operator, set equal to zero, and solve for the x which minimizes the derivative.

$$\nabla_x (E_\gamma(\hat{f}(x; x_0))) = \nabla f(x_0) + \frac{1}{2\gamma} (2x - 2x_0) = 0 \rightarrow$$

$$\nabla f(x_0) + \frac{1}{\gamma} (x - x_0) = 0 \rightarrow$$

$$x = x_0 - \gamma \nabla f(x_0)$$

The proximal operator is the x which minimizes the Moreau envelope, therefore:

$$\operatorname{prox}_\gamma = x_0 - \gamma \nabla f(x_0)$$

This has the form of a gradient descent update with step size γ . ■

Part B

Pseudo-code for the proximal gradient algorithm is as follows.

Initial Steps:

- Set initial beta value (such as a vector of zeros).
- Set initial objective function value, using $f(X,y,\lambda,\beta_0)$.
- Set up matrix to hold gradient values at each iteration.

Begin iterations at $i=2$.

Gradient Step:

- Update $gradient^{(i-1)}$ using gradient for $f(x)$ part of objective (see below).

Proximal Step:

- Calculate intermediate vector $z = \beta^{(i-1)} - \gamma * gradient^{(i-1)}$
- Update betas using $\beta^{(i)} = prox(z, \gamma, \lambda)$

Convergence Housekeeping Step:

- Update objective function using $\beta^{(i)}$.
- Check for convergence using abs change in objective function.

A few notes for implementing the proximal gradient algorithm for LASSO.

- LASSO: $\hat{\beta} = \underset{\beta}{\operatorname{argmin}} \{ \|y - X\beta\|_2^2 + \lambda \|\beta\|_1 \} = \underset{\beta}{\operatorname{argmin}} \{ f + g \}$
- Objective function is rewritten as: $(y - X\beta)^T(y - X\beta) + \lambda \sum |\beta|$
- Gradient of the non-penalty portion of the objective
function is $\nabla_{\beta}(f) = \nabla_{\beta} ((y - XB)^T(y - XB)) = -2X^T y + 2X^T X \beta$

Part B - Implementation in R

Implementation of the accelerated proximal gradient descent converged more quickly than the regular proximal gradient descent. See figure below for the objective functions.

```
> print(output$iter)
[1] 7687
> print(outputAccel$iter)
[1] 1853
5 > print(output$converged)
[1] 1
> print(outputAccel$converged)
[1] 1
```

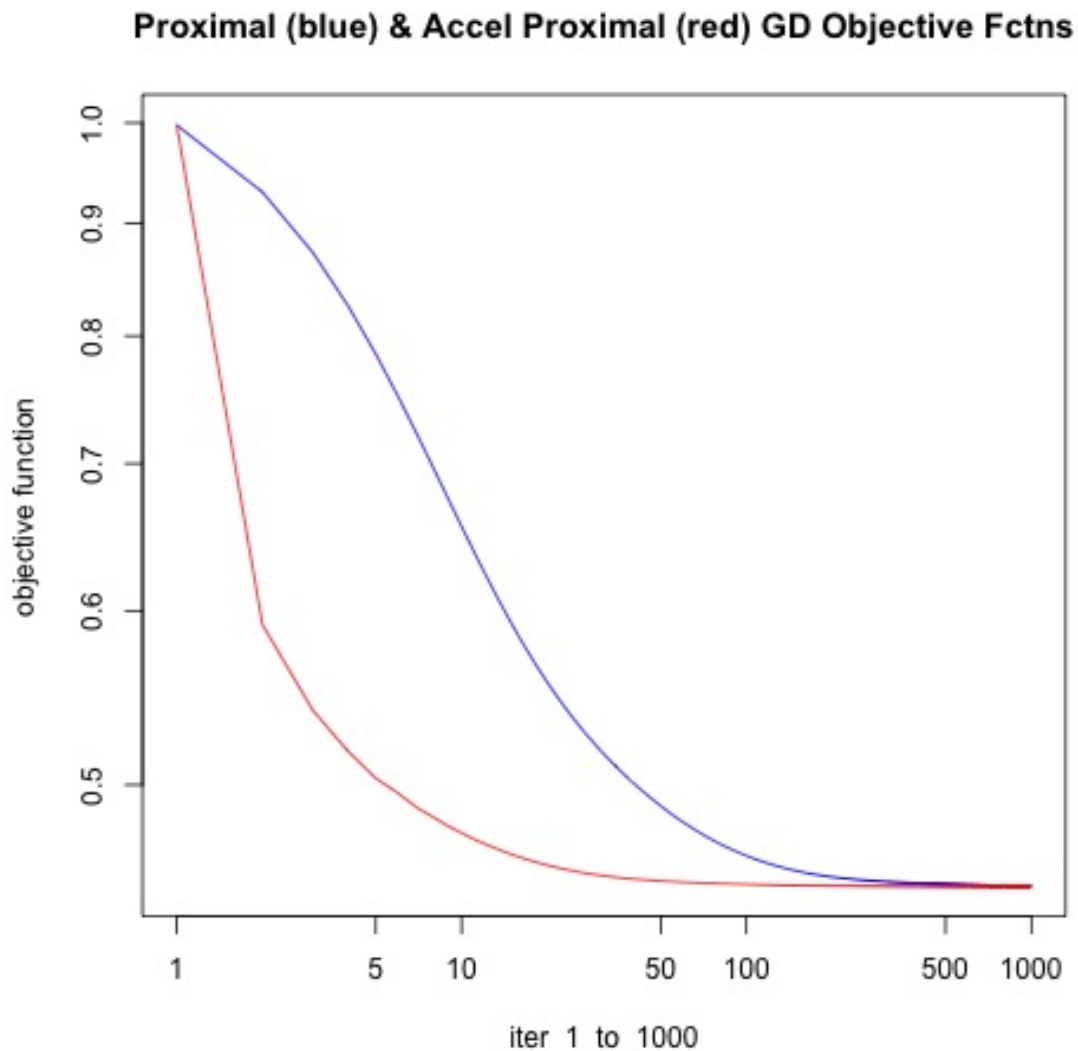


Figure 1: Prox & Accel Prox Objectives

For Accelerated Proximal, it did seem to descend consistently, but this was a surprise to me. In reading, I saw documentation of 'Nesterov Ripples', which noted that Accelerated Proximal Gradient Descent is not truly a descent algorithm.

Here is a comparison of the beta coefficient estimates I obtained with both methods, versus those calculated by glmnet. The coefficients look to be of the same magnitude and direction in both cases.

```

64 x 3 sparse Matrix of class "dgCMatrix"
      s0      proximal  accel.prox
age    0.0084994686  0.0171642800  0.01572673
sex   -0.1274718132 -0.1394743658 -0.13943695
5  bmi    0.3053578438  0.3051274718  0.30015079
map    0.1928804642  0.2009670773  0.20139367
tc     -0.0387865853 -0.0377155655 -0.06178009
ldl     .             -0.0376434751  .
hdl    -0.1635236545 -0.1426736417 -0.15577212
10  tch     .             0.0537160979  0.01436791
ltg    0.3230766076  0.3198859003  0.34286594
glu    0.0337240035  0.0373169619  0.03762443
age.2  0.0351045119  0.0427545290  0.04263404
bmi.2  0.0238787313  0.0231121365  0.02483871
15  map.2  .             -0.0015880181 -0.00208119
tc.2   .             .             .
ldl.2  .             .             .
hdl.2  .             .             .
tch.2  .             0.0182890006  0.03565516
20  ltg.2  .             .             0.01278624
glu.2  0.0636440798  0.0735089090  0.06831205
age.sex 0.0961729059  0.1030491523  0.10416798
age.bmi .             .             .
age.map 0.0145699540  0.0118708666  0.01113105
25  age.tc .             .             .
age.ldl -0.0389574124 -0.0517479493 -0.05308852
age.hdl 0.0258647293  0.0390369802  0.03935157
age.tch .             0.0005535911  0.00218560
age.ltg 0.0407916229  0.0539588720  0.05090272
30  age.glu 0.0184222925  0.0237785514  0.02458299
sex.bmi 0.0270461479  0.0366747723  0.03564408
sex.map 0.0387600107  0.0438562290  0.04318104
sex.tc  .             .             .
sex.ldl -0.0181954985 -0.0202717755 -0.01920683
35  sex.hdl 0.0429324533  0.0509889240  0.04841902
sex.tch .             -0.0103067180 -0.01352070
sex.ltg .             .             .
sex.glu .             0.0095186855  0.01010271
bmi.map 0.0837040826  0.0921837729  0.09113240
40  bmi.tc  -0.0002777938 -0.0127770877 -0.01684757
bmi.ldl .             .             .
bmi.hdl .             .             .
bmi.tch .             .             .

```


	bmi.ltg	.	.	.
45	bmi.glu	.	0.0015817727	0.00314029
	map.tc	0.0083855329	0.0139689484	0.02330329
	map.ldl	.	0.0088425035	.
	map.hdl	0.0233788843	0.0269141017	0.02262387
	map.tch	.	.	.
50	map.ltg	.	.	.
	map.glu	-0.0362793960	-0.0593616103	-0.05895615
	tc.ldl	.	0.0097864282	0.01510727
	tc.hdl	0.0072487602	0.0046025918	.
	tc.tch	-0.0477242186	-0.0772068542	-0.09664302
55	tc.ltg	-0.0098056312	-0.0369761727	-0.05132028
	tc.glu	.	.	.
	ldl.hdl	.	.	.
	ldl.tch	.	.	.
	ldl.ltg	0.0742412609	0.1157947622	0.13540195
60	ldl.glu	0.0067025050	0.0041732415	.
	hdl.tch	-0.0522309970	-0.0435042241	-0.04250106
	hdl.ltg	.	.	.
	hdl.glu	.	0.0235383572	0.03276363
	tch.ltg	-0.0607125224	-0.0786600477	-0.08710951
65	tch.glu	0.0214432611	0.0436959702	0.05905093
	ltg.glu	.	0.0023713767	0.00480166

Appendix: R Code

```
### SDS 385 - Exercises 06 - Proximal Gradient Descent for LASSO.

#Jennifer Starling
#7 October 2016

5 rm(list=ls()) #Clean workspace.

library(glmnet)
library(Matrix)

10 #Read in Diabetes.csv data.
X <- read.csv(file='/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Big_Data/
  Exercise 05 R Code/DiabetesX.csv',header=T)
y <- read.csv(file='/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Big_Data/
  Exercise 05 R Code/DiabetesY.csv',header=F)

15 #Scale X and y.
X = scale(X)
y = scale(y)

#-----
20 #LASSO objective function:
#Inputs:
# X = X matrix (scaled)
# y = response data (scaled)
# lambda = a chosen lambda value
25 # beta = a vector of beta coefficients.
#Output:
# Value of the LASSO objective function at specified inputs.
fx <- function(X,y,lambda,beta){
  f = (1/nrow(X)) * (t(y - X %*% beta) %*% (y - X %*% beta))
  g = lambda * sum(abs(beta))
30  obj = (f+g)
  return(as.numeric(obj))
}

35 #Test:
fx(X,y,lam,beta_glmnet)

#-----
#Proximal L1 Operator function: (soft thresholding operator)
40 #Inputs:
# x = vector of values.
# lambda = the scaling factor of the l1 norm.
# t = the step size.

45 #Output:
# Value of the soft-thresholding proximal operator.

prox_l1 <- function(x,gamma,tau=1) {

50  thresh <- gamma*tau
  prox = rep(0,length(x))

  idx.1 = which(x < -thresh)
  idx.2 = which(x > thresh)
```

```

55     idx.3 = which(abs(x) <= thresh)

    if (length(idx.1) > 0) prox[idx.1] = x[idx.1] + thresh
    if (length(idx.2) > 0) prox[idx.2] = x[idx.2] - thresh
    if (length(idx.3) > 0) prox[idx.3] = 0

60     return(prox)
}

#-----
65 #Gradient for differentiable (non-penalty) part of LASSO objective:
gradient <- function(X,y,beta){
    grad = (2/nrow(X)) * (t(X) %*% X %*% beta - t(X) %*% y )
    return(grad)
}

70 #-----
#Proximal Gradient Descent for L1 Norm Function:
#Inputs:
#   X = design matrix
75 #   y = response vector
#   gamma = step size
#   maxiter = maximum iterations
#   tol = tolerance for convergence
#   lambda = l1 norm penalty constant.
80 #Output:
#   List including estimated beta values and objective function.

proxGD <- function(X,Y,gamma=.01,maxiter=50,tol=1E-10,lambda=.1){

85     i=0                #Initialize iterator.
    converged <- 0       #Indicator for whether convergence met.

    #1. Initialize matrix to hold beta vector for each iteration.
    betas <- matrix(0,nrow=maxiter,ncol=ncol(X))
90     betas[1,] <- rep(0,ncol(X)) #Initialize beta vector to 0 to start.

    #2. Initialize values for objective function.
    obj <- rep(0,maxiter) #Initialize vector to hold loglikelihood fctn.
    obj[1] <- fx(X,y,lambda,betas[1,])

95     #3. Initialize matrix to hold gradients for each iteration.
    grad <- matrix(0,nrow=maxiter,ncol=ncol(X))

    for (i in 2:maxiter){
100         #STEP 1: Gradient Step.

        #Calc gradient.
        #grad[i-1,] = (2/nrow(X)) * (t(X) %*% X %*% betas[i-1,] - t(X) %*% y )
        grad[i-1,] = gradient(X,y,betas[i-1,])

105         #Determine intermediate point.
        z = betas[i-1,] - gamma*grad[i-1,]

        #STEP 2: Proximal step.
110         betas[i,] = prox_l1(z,gamma,tau=lambda)

        #Update objective function.
        obj[i] = fx(X,y,lambda=lambda,beta=betas[i,])
    }
}

```

```

115     #Check if convergence met: If yes, exit loop.
        if (abs(obj[i]-obj[i-1])/abs(obj[i-1]+1E-3) < tol ){
            converged=1;
            break;
        }
120 } #end for loop

        return(list(obj=obj, betas=betas, beta_hat=betas[i,], converged=converged,
            iter=i))
    } #end function

125 #-----
#Accelerated Proximal Gradient Descent for L1 Norm Function:
#(Nesterov)
#Inputs:
#   X = design matrix
#   y = response vector
130 #   gamma = step size
#   maxiter = maximum iterations
#   tol = tolerance for convergence
#   lambda = l1 norm penalty constant.
135 #Output:
#   List including estimated beta values and objective function.

accelProxGD <- function(X,Y,gamma=.01,maxiter=50,tol=1E-10,lambda=.1){

140     i=0                #Initialize iterator.
    converged <- 0        #Indicator for whether convergence met.

    #1. Initialize matrix to hold beta vector for each iteration.
    betas <- matrix(0,nrow=maxiter,ncol=ncol(X))
145     betas[1,] <- rep(0,ncol(X)) #Initialize beta vector to 0 to start.

    #2. Initialize values for objective function.
    obj <- rep(0,maxiter) #Initialize vector to hold loglikelihood fctn.
    obj[1] <- fx(X,y,lambda,betas[1,])

150     #3. Initialize matrix to hold gradients for each iteration.
    grad <- matrix(0,nrow=maxiter,ncol=ncol(X))
    grad[1,] = gradient(X,y,betas[1,])

155     #4. Initialize vectors to hold Nesterov update values.
    z = matrix(0,nrow=maxiter,ncol=ncol(X))
    s = rep(0,maxiter)

    #Set up first z value. (Used a regular gradient calculation for beta0.)
160     #z[1,] = betas[1,] - gamma * grad[1,]

    #Set up scalar s terms. Ok before main loop, as do not depend on other terms'
    #updates.
    for (j in 2:maxiter){
        s[j] = (1 + sqrt(1 + 4*(s[j-1])^2)) / 2
165     }

    #Loop through iterations until converged or maxiter met.
    for (i in 2:maxiter){

170         #STEP 1: Gradient Step.

```

```

#Calc gradient.
grad[i-1,] = gradient(X,y,z[i-1,])

#Update intermediate u term.
u = z[i-1,] - gamma * grad[i-1,]

#STEP 2: Proximal step; update betas.
betas[i,] = prox_l1(u,gamma,tau=lambda)

#STEP 3: Nesterov step; update Nesterov momentum z.
z[i,] = betas[i-1,] + ((s[i-1]-1)/s[i]) * (betas[i-1,] - betas[i,])

#Update objective function.
obj[i] = fx(X,y,lambda=lambda,beta=betas[i,])

#Check if convergence met: If yes, exit loop.
#if (abs(obj[i]-obj[i-1])/abs(obj[i-1]+1E-10) < tol ){
#   converged=1;
#   break;
#}
} #end for loop

return(list(obj=obj, betas=betas, beta_hat=betas[i,], converged=converged,
            iter=i,s=s))
} #end function

#-----

#Run proximal gradient descent & accelerated proximal gradient descent.
lam=.01
output <- proxGD(X,y,gamma=.01,maxiter=1000,tol=1E-10,lambda=lam)
outputAccel <- accelProxGD(X,y,gamma=.01,maxiter=1000,tol=1E-10,lambda=lam)

#Iterations to convergence:
print(output$iter)
print(outputAccel$iter)
print(output$converged)
print(outputAccel$converged)

#Compare results to glmnet:
myLasso <- glmnet(X,y,family='gaussian',lambda=lam) #fit glmnet model.
beta_glmnet <- myLasso$beta #Save glmnet betas.
cbind(glmnet=beta_glmnet,
      proximal=output$beta_hat,
      accel.prox=round(outputAccel$beta_hat,8)) #output comparison

#Plot objective function.
plot(1:output$iter,output$obj[1:output$iter],type='l',log='xy',col='blue',xlab=
      paste('iter ',1,' to ',output$iter),
      ylab='objective function')
lines(1:outputAccel$iter,outputAccel$obj[1:outputAccel$iter],type='l',col='red',
      xlab=paste('iter ',1,' to ',outputAccel$iter),
      ylab='objective function')

#Plot convergence of betas.
idx = which(output$beta_hat>0)

```

```
#idxPlot = sample(idx,9,replace=F)
for (j in idx){
  plot(1:length(output$betas[,j]),output$betas[,j],xlab='iter',ylab=paste('
    betahat',j),type='l',col='blue')
230   abline(h=beta_glmnet[j],col='red')
}

235 #-----
#Run accelerated proximal gradient descent.
```