

**SDS 385 Ex 07:**  
**ADMM for LASSO**

October 18, 2016

**Jennifer Starling**

## ADMM for LASSO

### Algorithm Details

I used the ADMM algorithm as detailed in Section 6.4 of Boyd, "Distributed Optimization and Statistical Learning via the Alternating Direction Method of Multipliers", 2011.

The setting for this algorithm as applied to lasso is below.

#### Objective Function:

minimize  $\frac{1}{2} \|X\beta - y\|_2^2 + \lambda \|x\|_1$   
with  $\lambda > 0$  as the l1 regularization penalty parameter.

#### ADMM Form of Problem:

minimize  $f(x) + g(x)$   
subject to  $x - z = 0$   
where:  
 $f(x) = \frac{1}{2} \|X\beta - y\|_2^2$   
 $g(x) = \lambda \|x\|_1$

#### ADMM Steps:

$$\begin{aligned}\beta^{k+1} &= (X^T X + \rho I)^{-1} (X^T y + \rho(z^k - u_k)) \\ z^{k+1} &= S_{\lambda/\rho}(\beta^{k+1} + u^k) \\ u^{k+1} &= u^k + \beta^{k+1} - z^{k+1}\end{aligned}$$

where  $S_{\lambda/\rho}(a) = \text{sign}(a)(|a| - \lambda/\rho)_+$   
and  $u^k$  is the augmented lagrangian at the  $k^{th}$  iteration.

#### Implementation Notes:

- Since the step size is fixed, to increase efficiency, I cached the matrix inverse  $(X^T X + \rho I)^{-1}$  before beginning iterations.
- Boyd notes that when not using fixed step size, it is best to still cache this inverse, and only update it when the  $\rho$  step size is updated.

## My Implementation

My ADMM implementation for LASSO converged in 230 iterations. This was considerably faster than my proximal gradient descent (7687 iterations) and my accelerated proximal gradient descent (1853 iterations). I used  $\lambda = .01$  for the l1 lasso penalty and  $\rho = .01$  for the step size. I used the convergence criteria of absolute change in the lasso objective function  $< tol$ , where  $tol = 1E - 10$ .

Note that the lambda and rho were the same values used for exercise 6, for consistency. In practice, I would select an optimal lambda value via cross-validation.

The plot of the lasso objective function is as follows.




Figure 1: LASSO objective using ADMM( $\hat{\alpha}$ )

The estimated beta coefficients obtained using admm are compared with the coefficients from glmnet below. Most appear to be of reasonably similar sign and magnitude.

	glmnet	admm
age	0.0084994686	0.03007395
sex	-0.1274718132	-0.16390462
bmi	0.3053578438	0.28268400
map	0.1928804642	0.21180397
tc	-0.0387865853	-0.10040001
ldl	.	.
hdl	-0.1635236545	-0.10958174
tch	.	0.05421951
ltg	0.3230766076	0.42695740
glu	0.0337240035	0.03968565
age.2	0.0351045119	0.04230375
bmi.2	0.0238787313	0.02913876
map.2	.	-0.00496873
tc.2	.	3.32992803

Jennifer Starling

	ldl.2	.	1.78879470
	hdl.2	.	0.86057254
	tch.2	.	0.45016330
	ltg.2	.	0.71081295
20	glu.2	0.0636440798	0.07115055
	age.sex	0.0961729059	0.09425725
	age.bmi	.	-0.00974232
	age.map	0.0145699540	0.01097956
	age.tc	.	-0.10307316
25	age.ldl	-0.0389574124	-0.03850246
	age.hdl	0.0258647293	0.13295890
	age.tch	.	0.11473838
	age.ltg	0.0407916229	0.07678043
	age.glu	0.0184222925	0.03852093
30	sex.bmi	0.0270461479	0.04070618
	sex.map	0.0387600107	0.05455798
	sex.tc	.	0.25201571
	sex.ldl	-0.0181954985	-0.20580754
	sex.hdl	0.0429324533	-0.06940488
35	sex.tch	.	-0.07673176
	sex.ltg	.	-0.07203117
	sex.glu	.	0.02857881
	bmi.map	0.0837040826	0.09607870
	bmi.tc	-0.0002777938	-0.22702825
40	bmi.ldl	.	0.18330034
	bmi.hdl	.	0.08852731
	bmi.tch	.	-0.01969956
	bmi.ltg	.	0.07900017
	bmi.glu	.	0.01519782
45	map.tc	0.0083855329	0.28434387
	map.ldl	.	-0.19175559
	map.hdl	0.0233788843	-0.11053647
	map.tch	.	-0.03567046
	map.ltg	.	-0.09184804
50	map.glu	-0.0362793960	-0.08424349
	tc.ldl	.	-4.61510706
	tc.hdl	0.0072487602	-1.93720146
	tc.tch	-0.0477242186	-1.10984540
	tc.ltg	-0.0098056312	-2.26554521
55	tc.glu	.	-0.08709231
	ldl.hdl	.	1.26150669
	ldl.tch	.	0.52479375
	ldl.ltg	0.0742412609	1.70305929
	ldl.glu	0.0067025050	0.03256134
60	hdl.tch	-0.0522309970	0.62815286
	hdl.ltg	.	0.84045023
	hdl.glu	.	0.12679103
	tch.ltg	-0.0607125224	0.15900559
	tch.glu	0.0214432611	0.14614305

65

<code>ltg.glu</code> .	0.04353096
------------------------	------------

## Appendix: R Code

```
### SDS 385 - Exercises 07 - ADMM for LASSO.
### ADMM = Alternating Direction Method of Multipliers

#Jennifer Starling
5 #18 October 2016

rm(list=ls()) #Clean workspace.

library(glmnet)
10 library(Matrix)

#Read in Diabetes.csv data.
X <- read.csv(file='/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Big_Data/
  Exercise 05 R Code/DiabetesX.csv',header=T)
y <- read.csv(file='/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Big_Data/
  Exercise 05 R Code/DiabetesY.csv',header=F)
15

#Scale X and y.
X = scale(X)
y = scale(y)

#-----
#LASSO objective function:
#Inputs:
# X = X matrix (scaled)
# y = response data (scaled)
25 # lambda = a chosen lambda value
# beta = a vector of beta coefficients.
#Output:
# Value of the LASSO objective function at specified inputs.
fx <- function(X,y,lambda,beta){
30   obj = (1/2) * sum((y - X %*% beta) ^ 2) + lambda * sum(abs(beta))
   return(obj)
}

#-----
35 #Proximal L1 Operator function: (soft thresholding operator)
prox_l1 <- function(x, lambda){

  # Computes the soft thresholding estimator
  # -----
40 # Args:
  # - x: vector of the observations
  # - lambda: penalization parameter (threshold)
  # Returns:
  # - theta: the soft thresholding estimator
45 # -----
  theta <- sign(x) * pmax(rep(0, length(x)), abs(x) - lambda)
  return (theta)
}

50 #-----
#ADMM for Lasso:
#Inputs:
# X = design matrix (A)
# y = response vector (b)
```

```

55 # rho = step size
# maxiter = maximum iterations
# eps_abs = primal tolerance for convergence (epsilon_abs)
# eps_rel = dual tolerance for convergence (epsilon_rel)
# lambda = l1 norm penalty constant.
60 #Output:
# List including estimated beta (x) values and objective function.
#Note: In optimization notation, A=X, b=Y, x=beta (minimizing x).

admmLasso = function(X,Y,rho=1,lambda=.1,maxiter=1000,e_abs=1E-3,e_rel=1E-6){
65
#Define dimensions n and p.
n = nrow(X)
p = ncol(X)

70 #Rescale lambda to match glmnet results.
lambda = lambda * n

#Define function Euclidian (l2) norm of a vector.
l2norm <- function(x) sqrt(sum(x^2))

75
i=0 #Initialize iterator.
converged <- 0 #Indicator for whether convergence met.

#Initialize data structures.
80 betas <- matrix(0,nrow=maxiter,ncol=p) #holds beta vector for each iteration.
obj <- rep(0,maxiter) #Initialize vector to hold loglikelihood fctn.
z = matrix(0,nrow=maxiter,ncol=p) #Initialize z vector to all zeros.

#Initialize values.
85 obj[1] <- fx(X,y,lambda,betas[1,]) #Initialize objective.
betas[1,] <- rep(0,p) #Initialize beta vector to 0 to start.
u = rep(0,p) #Initialize the lagrangian to all zeros.

#Pre-cache matrix inverse and Xty, since using fixed step size for each iter.
90 Xty = crossprod(X,y)
inv = solve(crossprod(X) + diag(rep(rho,p)))

#Initialize residual vectors.
s = 0 #dual residual
95 r = 0 #primal residual

#ADMM looping.
for (i in 2:maxiter){
100
#Update betas.
betas[i,] = inv %*% (Xty + rho * (z[i,]-u) )

#Update z.
z[i,] = prox_l1(betas[i,] + u,lambda/rho)
105
#Update u (lagrangian).
u = u + betas[i,] - z[i,]

#Update objective function.
110 obj[i] = fx(X,y,lambda=lambda,beta=betas[i,])

#-----
#Convergence check:

```

```

115     #Calculate residuals for iteration i.
        r = betas[i,] - z[i,]
        s = -rho * (z[i,] - z[i-1,])

        r.norm = l2norm(r)
120     s.norm = l2norm(s)

        e.primal = sqrt(p)*e_abs + e_rel * max(l2norm(betas[i,]), l2norm(z[i,]))
        e.dual = sqrt(p)*e_abs + e_rel * l2norm(u)

125     if (r.norm <= e.primal && s.norm <= e.dual){
        converged=1
        break
    }
    #-----
130 }

    #Return function values.
    return(list(obj=obj, betas=betas, beta_hat=betas[i,], converged=converged,
        iter=i))
}
135 #-----

#Run admm for lasso.
output <- admmLasso(X,y,rho=5,lambda=.01,maxiter=1000,e_abs=1E-6,e_rel=1E-2)

140 #Iterations to convergence:
print(output$iter)
print(output$converged)

#Plot objective function.
145 #jpeg(file='/Users/jennstarling/UTAustin/2016_Fall_SDS_385_Big_Data/Exercise_07
    LaTeX_Files/admm_objective.jpg')
plot(1:output$iter,output$obj[1:output$iter],type='l',col='blue',log='xy',
    main='Lasso Objective Function',xlab='iter',ylab='objective')
#dev.off()

150 #Compare results to glmnet:
myLasso <- glmnet(X,y,family='gaussian',alpha=1,lambda=.01,intercept=F,standardize
    =F) #Fit lasso glmnet model.
beta_glmnet <- myLasso$beta #Save glmnet betas.
cbind(glmnet=beta_glmnet,admm=round(output$beta_hat,8)) #Output comparison

```